

Linux Basic Commands

- [1. ls \(List Directory Contents\)](#)
- [2. cd \(Change Directory\)](#)
- [3. mv \(Move/Rename Files or Directories\)](#)
- [4. cp \(Copy Files or Directories\)](#)
- [5. rm \(Remove Files or Directories\)](#)
- [6. touch \(Create Empty Files or Update Timestamps\)](#)
- [7. mkdir \(Make Directories\)](#)

[Practical Tips](#)

[Practice Tasks](#)

Linux File Permissions

- [1. chmod \(Change File Mode\)](#)
- [2. chown \(Change Owner\)](#)
- [3. umask \(User Mask\)](#)
- [4. stat \(Display File Status\)](#)

[Practical Tips](#)

[Practice Tasks](#)

Linux File Viewing Tools

- [1. cat \(Concatenate and Display Files\)](#)
- [2. less \(View File Contents Interactively\)](#)
- [3. more \(View File Contents Page by Page\)](#)
- [4. head \(Display First Lines of a File\)](#)
- [5. tail \(Display Last Lines of a File\)](#)
- [6. grep \(Search Text Using Patterns\)](#)

[Practical Tips](#)

[Practice Tasks](#)

Linux Process Management (Including PM2)

- [1. ps \(Process Status\)](#)
- [2. top \(Display System Processes\)](#)
- [3. htop \(Interactive Process Viewer\)](#)
- [4. kill \(Terminate Processes\)](#)
- [5. nice \(Set Process Priority\)](#)
- [6. renice \(Change Running Process Priority\)](#)
- [7. fg \(Foreground a Process\)](#)
- [8. bg \(Background a Process\)](#)
- [9. pm2 \(Process Manager for Node.js Applications\)](#)

[Practical Tips](#)

Practice Tasks

Linux Package Management

1. apt (Advanced Package Tool)
2. yum (Yellowdog Updater, Modified)
3. dnf (Dandified YUM)
4. dpkg (Debian Package)
5. rpm (RPM Package Manager)

Practical Tips

Practice Tasks

Linux Network Utilities

1. ping (Packet Internet Groper)
2. netstat (Network Statistics)
3. ss (Socket Statistics)
4. ip (IP Configuration)
5. nmcli (NetworkManager Command Line Interface)
6. curl (Client URL)
7. wget (Web Get)

Comparison: curl vs. wget

Practical Tips

Practice Tasks

Linux DNS and Hosts File

1. /etc/hosts (Local Hosts File)
2. /etc/resolv.conf (Resolver Configuration File)
3. DNS Configuration Basics

Practical Tips

Practice Tasks

2. /etc/resolv.conf (Resolver Configuration File)
3. DNS Configuration Basics

Practical Tips

Practice Tasks

Linux Cron Jobs and Scheduling

1. crontab (Cron Table)
2. at (One-Time Task Scheduling)
3. Automating Tasks

Practical Tips

Practice Tasks

Linux Systemd and Services

1. systemctl (System Control)
2. service (Service Management)
3. journalctl (Systemd Journal Logs)

Practical Tips

[Practice Tasks](#)

[Linux User Management](#)

- [1. adduser \(Add a User\)](#)
- [2. usermod \(Modify a User\)](#)
- [3. passwd \(Change User Password\)](#)
- [4. groups \(Display User Groups\)](#)

[Practical Tips](#)

[Practice Tasks](#)

[Linux Group Management](#)

- [1. groupadd \(Add a Group\)](#)
- [2. gpasswd \(Administer Group Password and Membership\)](#)
- [3. usermod -aG \(Add User to Groups\)](#)

[Practical Tips](#)

[Practice Tasks](#)

[Linux File System Basics](#)

- [1. df \(Disk Free\)](#)
- [2. du \(Disk Usage\)](#)
- [3. lsblk \(List Block Devices\)](#)
- [4. fdisk \(Partition Table Manipulator\)](#)
- [5. mkfs \(Make File System\)](#)

[Practical Tips](#)

[Practice Tasks](#)

[Linux Mounting Drives \(Detailed Explanation\)](#)

- [1. mount \(Mount a File System\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

- [2. umount \(Unmount a File System\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

- [3. /etc/fstab \(File System Table\)](#)

[Overview](#)

[File Format](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[Practical Tips](#)

[Practice Tasks](#)

[Linux Basic LVM Concepts](#)

[Overview of LVM](#)

[1. pvcreate \(Create Physical Volume\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[2. vgcreate \(Create Volume Group\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[3. lvcreate \(Create Logical Volume\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[Practical Example: Setting Up an LVM Volume](#)

[Scenario](#)

[Steps](#)

[Practical Tips](#)

[Practice Tasks](#)

[Linux File Permissions](#)

[Overview of File Permissions](#)

[1. chmod \(Change Mode\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[2. chown \(Change Owner\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[3. setfacl \(Set File Access Control Lists\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[4. getfacl \(Get File Access Control Lists\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[Practical Tips](#)

[Practice Tasks](#)

[Linux Logs and Log Management](#)

[Overview of Logs](#)

[1. tail \(View End of File\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[2. grep \(Search Text in Files\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[3. /var/log \(System Log Directory\)](#)

[Overview](#)

[Common Log Files](#)

[Detailed Examples](#)

[Troubleshooting](#)

[4. logrotate \(Rotate Log Files\)](#)

[Overview](#)

[Configuration File](#)

[Detailed Examples](#)

[Troubleshooting](#)

[Practical Tips](#)

[Practice Tasks](#)

[Linux Process Management](#)

[Overview of Process Management](#)

[1. ps \(Process Status\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[2. top \(Table of Processes\)](#)

[Overview](#)

[Common Keys \(Interactive\)](#)

[Detailed Examples](#)

[Troubleshooting](#)

[3. htop \(Enhanced top\)](#)

[Overview](#)

[Common Keys \(Interactive\)](#)

[Detailed Examples](#)

[Troubleshooting](#)

[4. kill \(Send Signal to Process\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[5. nice \(Set Process Priority\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[Practical Tips](#)

[Practice Tasks](#)

[Linux Package Management](#)

[Overview of Package Management](#)

[1. apt \(Advanced Package Tool\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[2. yum/dnf \(Yellowdog Updater Modified/Dandified YUM\)](#)

[Overview](#)

[Common Options \(yum/dnf\)](#)

[Detailed Examples](#)

[Troubleshooting](#)

[3. dpkg \(Debian Package Manager\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[4. rpm \(RPM Package Manager\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[Practical Tips](#)

[Practice Tasks](#)

[Linux Firewall Management](#)

[Overview of Firewall Management](#)

[1. ufw \(Uncomplicated Firewall\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[2. iptables \(IP Tables\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[3. firewalld \(Firewall Daemon\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[Practical Tips](#)

[Linux SSH and Remote Access](#)

[Overview of SSH and Remote Access](#)

[1. ssh \(Secure Shell Client\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[2. scp \(Secure Copy\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[3. sshd_config \(SSH Daemon Configuration\)](#)

[Overview](#)

[Common Directives](#)

[Detailed Examples](#)

[Troubleshooting](#)

[4. Key-Based Authentication](#)

[Overview](#)

[Steps to Set Up](#)

[Detailed Example](#)

[Troubleshooting](#)

[Practical Tips](#)

[Practice Tasks](#)

[Linux System Monitoring and Performance](#)

[Overview of System Monitoring](#)

[1. sar \(System Activity Reporter\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[2. vmstat \(Virtual Memory Statistics\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[3. iostat \(Input/Output Statistics\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[4. free \(Memory Usage\)](#)

[Overview](#)

[Common Options](#)

[Detailed Examples](#)

[Troubleshooting](#)

[Practical Tips](#)

[Practice Tasks](#)

Linux Basic Commands

This document covers the foundational Linux commands for navigating and managing files and directories: `ls`, `cd`, `mv`, `cp`, `rm`, `touch`, and `mkdir`. Each command includes a description, common options, and practical examples.

1. ls (List Directory Contents)

- **Description:** Displays files and directories in the current directory.
- **Common Options:**
 - `-l`: Long format (shows permissions, owner, size, etc.).
 - `-a`: Includes hidden files (those starting with `.`).
 - `-h`: Human-readable file sizes (used with `-l`).
 - `-t`: Sort by modification time (newest first).
 - `-r`: Reverse sort order.

Examples:

```
ls          # List files and directories
ls -l       # List in long format
ls -la      # List all files, including hidden, in long format
ls -lh      # List with human-readable sizes
ls -ltr     # List in long format, sorted by time, reversed
```

2. cd (Change Directory)

- **Description:** Navigates to a specified directory.
- **Common Options:** None typically used, but special arguments exist:
 - `..`: Move up one directory.
 - `~`: Go to the user's home directory.
 - `-`: Return to the previous directory.

Examples:

```
cd /home/user/docs # Navigate to the docs directory
cd ..              # Move up one directory
cd ~               # Go to home directory
cd -               # Return to the previous directory
cd                # Go to home directory (no argument)
```

3. mv (Move/Rename Files or Directories)

- **Description:** Moves or renames files and directories.
- **Common Options:**
 - `-i`: Prompt before overwriting.
 - `-f`: Force overwrite without prompting.
 - `-v`: Verbose, show what is being done.

Examples:

```
mv file.txt /home/user/docs/      # Move file.txt to docs directory
mv file.txt newname.txt           # Rename file.txt to newname.txt
mv -i file.txt /home/user/docs/   # Move with overwrite prompt
mv -v myfolder /home/user/backup/ # Move directory with verbose output
```

4. cp (Copy Files or Directories)

- **Description:** Copies files or directories.
- **Common Options:**
 - `-r`: Recursive, copy directories and their contents.
 - `-i`: Prompt before overwriting.
 - `-v`: Verbose, show what is being done.

Examples:

```
cp file.txt /home/user/backup/     # Copy file.txt to backup directory
cp -r myfolder /home/user/backup/  # Copy myfolder and its contents
cp -i file.txt /home/user/backup/  # Copy with overwrite prompt
cp -v file.txt file_copy.txt       # Copy with verbose output
```

5. rm (Remove Files or Directories)

- **Description:** Deletes files or directories.
- **Common Options:**
 - `-r`: Recursive, delete directories and their contents.
 - `-f`: Force deletion without prompting.
 - `-i`: Prompt before each deletion.
 - `-v`: Verbose, show what is being deleted.

Examples:

```
rm file.txt          # Delete file.txt
rm -r myfolder       # Delete myfolder and its contents
rm -rf myfolder      # Force delete myfolder without prompting
rm -i file.txt       # Delete with prompt
rm -v *.txt          # Delete all .txt files with verbose output
```

- **Caution:** Use `rm -rf` carefully, as it deletes files permanently without recovery.

6. touch (Create Empty Files or Update Timestamps)

- **Description:** Creates empty files or updates the timestamp of existing files.
- **Common Options:**
 - `-m`: Update modification time only.
 - `-a`: Update access time only.
 - `-c`: Do not create the file if it doesn't exist.

Examples:

```
touch newfile.txt    # Create empty file newfile.txt
touch -m existing.txt # Update modification time of existing.txt
touch -c nonexistent.txt # Do nothing if file doesn't exist
touch file1.txt file2.txt # Create multiple files
```

7. mkdir (Make Directories)

- **Description:** Creates new directories.
- **Common Options:**
 - `-p`: Create parent directories as needed (no error if they exist).
 - `-v`: Verbose, show each directory created.

Examples:

```
mkdir myfolder        # Create a directory named myfolder
mkdir -p /home/user/docs/new # Create nested directories
mkdir -v folder1 folder2 # Create multiple directories with verbose output
```

Practical Tips

- **Tab Completion:** Press `Tab` to auto-complete file or directory names.
- **Command History:** Use `up` and `down` arrows to navigate previous commands.
- **Combine Commands:** Use `&&` to run commands sequentially (e.g., `mkdir myfolder && cd myfolder`).
- **Safety with `rm`:** Double-check paths with `rm -rf` to avoid accidental deletion.
- **Practice Environment:** Use a virtual machine or non-critical directory for testing.

Practice Tasks

1. Create a directory `test`, navigate into it, and create an empty file `example.txt`.
2. Copy `example.txt` to `backup.txt` in the same directory.
3. Rename `backup.txt` to `backup_copy.txt`.
4. List all files in long format, including hidden ones.
5. Delete `example.txt` and the `test` directory.

Solution:

```
mkdir test
cd test
touch example.txt
cp example.txt backup.txt
mv backup.txt backup_copy.txt
ls -la
rm example.txt
cd ..
rm -r test
```

Linux File Permissions

This document covers Linux file permission commands: `chmod`, `chown`, `umask`, and `stat`. Each command includes a description, common options, and practical examples to help you understand and manage file permissions effectively.

1. chmod (Change File Mode)

- **Description:** Modifies file or directory permissions (read, write, execute) for owner, group, and others.
- **Permission Types:**
 - `r` (read): 4
 - `w` (write): 2
 - `x` (execute): 1
- **Permission Levels:**
 - `u`: User (owner)
 - `g`: Group
 - `o`: Others
 - `a`: All (user, group, others)
- **Common Options:**
 - `-R`: Recursive, apply to directories and their contents.
 - `-v`: Verbose, show changes made.
- **Syntax:**
 - Numeric mode: `chmod [permissions] file` (e.g., `755 = rwxr-xr-x`).
 - Symbolic mode: `chmod [who][+|-|=][permissions] file` (e.g., `u+x` adds execute for owner).

Examples:

```
chmod 644 file.txt      # Set owner: read/write (6), group/others: read (4)
chmod 755 script.sh     # Set owner: rwx (7), group/others: rx (5)
chmod -R 700 myfolder   # Set owner: rwx, group/others: none, recursively
chmod u+x script.sh     # Add execute permission for owner
chmod go-r file.txt     # Remove read permission for group and others
chmod -v a+w file.txt   # Add write for all, with verbose output
```

2. chown (Change Owner)

- **Description:** Changes the owner and/or group of a file or directory.
- **Common Options:**

- **-R**: Recursive, apply to directories and their contents.
- **-v**: Verbose, show changes made.
- **-h**: Modify symbolic link itself, not the target.
- **Syntax:**
 - `chown [user]:[group] file`
 - Omit `:[group]` to change only the owner.

Examples:

```
chown user1 file.txt          # Change owner to user1
chown user1:group1 file.txt   # Change owner to user1 and group to group1
chown -R user1:group1 myfolder # Change owner and group recursively
chown -v user2 file.txt       # Change owner with verbose output
chown -h user1 symlink        # Change owner of symlink itself
```

3. umask (User Mask)

- **Description:** Sets the default permissions for newly created files and directories by subtracting from the default permissions (usually `666` for files, `777` for directories).
- **Common Values:**
 - `022`: Files get `644` (rw-r--r--), directories get `755` (rwxr-xr-x).
 - `002`: Files get `664` (rw-rw-r--), directories get `775` (rwxrwxr-x).
- **Syntax:**
 - `umask [value]`: Set the umask.
 - `umask`: Display current umask.

Examples:

```
umask 022          # Set umask to 022 (default for most systems)
umask              # Display current umask (e.g., 0022)
umask 002          # Set umask to 002 (more permissive for group)
touch newfile.txt   # Create file with permissions 664 (if umask is 002)
mkdir newdir        # Create directory with permissions 775 (if umask is 002)
```

4. stat (Display File Status)

- **Description:** Displays detailed file or directory information, including permissions, ownership, timestamps, and more.
- **Common Options:**
 - **-f**: Display filesystem status instead of file status.

- `--format`: Customize output format.

Examples:

```
stat file.txt          # Display detailed info about file.txt
stat -f /home          # Display filesystem info for /home
stat --format="%n %U %G %A" file.txt # Show name, owner, group, and permissions
```

Practical Tips

- **Understanding Permissions:**
 - Numeric: Sum permissions for owner, group, others (e.g., `755` = `rw` for owner, `r-x` for group/others).
 - Symbolic: Use `+` to add, `-` to remove, `=` to set exact permissions.
- **Check Permissions:** Use `ls -l` to view permissions before and after using `chmod` or `chown`.
- **Default umask:** Common default is `022`. Adjust in `~/.bashrc` or `/etc/profile` for persistence.
- **stat for Debugging:** Use `stat` to verify permissions, ownership, or timestamps when troubleshooting.
- **Safety:** Use `-i` with `chmod` or `chown -R` to avoid unintended changes.

Practice Tasks

1. Create a file `test.txt` and set its permissions to `rw-r--r--` (644).
2. Change the owner of `test.txt` to `user1` and group to `group1`.
3. Set umask to `002`, then create a file `newfile.txt` and directory `newdir`. Check their permissions.
4. Use `stat` to display the owner, group, and permissions of `test.txt`.
5. Recursively set permissions of a directory `myfolder` to `700`.

Solution:

```
touch test.txt
chmod 644 test.txt
chown user1:group1 test.txt
umask 002
touch newfile.txt
mkdir newdir
ls -l newfile.txt newdir # Should show 664 for file, 775 for directory
stat --format="%n %U %G %A" test.txt
```

```
mkdir myfolder
chmod -R 700 myfolder
```

Linux File Viewing Tools

This document covers Linux commands for viewing and searching file contents: `cat`, `less`, `more`, `head`, `tail`, and `grep`. Each command includes a description, common options, and practical examples to help you effectively view and manipulate text files.

1. cat (Concatenate and Display Files)

- **Description:** Displays the contents of one or more files, concatenates files, or redirects output to a new file.
- **Common Options:**
 - `-n`: Number all output lines.
 - `-b`: Number non-blank lines.
 - `-s`: Squeeze multiple blank lines into one.

Examples:

```
cat file.txt           # Display contents of file.txt
cat -n file.txt        # Display file.txt with line numbers
cat file1.txt file2.txt # Concatenate and display file1.txt and file2.txt
cat file1.txt file2.txt > combined.txt # Concatenate files into combined.txt
cat -s file.txt        # Display file.txt with squeezed blank lines
```

2. less (View File Contents Interactively)

- **Description:** Displays file contents one page at a time, allowing navigation (scrolling, searching).
- **Common Options:**
 - `-N`: Show line numbers.
 - `-i`: Ignore case when searching.
 - `-S`: Chop long lines instead of wrapping.
- **Navigation:**
 - `Space`: Next page.
 - `b`: Previous page.
 - `/pattern`: Search for `pattern`.

- **q**: Quit.

Examples:

```
less file.txt          # View file.txt interactively
less -N file.txt       # View with line numbers
less -i file.txt       # View with case-insensitive search
less -S longlines.txt  # View without line wrapping
```

3. more (View File Contents Page by Page)

- **Description:** Displays file contents one page at a time, with less interactivity than **less**.
- **Common Options:**
 - **+n**: Start at line number **n**.
 - **-n**: Display **n** lines per page.
 - **+ /pattern**: Start at the first occurrence of **pattern**.
- **Navigation:**
 - **Space**: Next page.
 - **Enter**: Next line.
 - **q**: Quit.

Examples:

```
more file.txt          # View file.txt page by page
more +10 file.txt       # Start at line 10
more -5 file.txt        # Display 5 lines per page
more +/error log.txt    # Start at first occurrence of "error"
```

4. head (Display First Lines of a File)

- **Description:** Displays the first 10 lines of a file by default.
- **Common Options:**
 - **-n N**: Display the first **N** lines.
 - **-c N**: Display the first **N** bytes.

Examples:

```
head file.txt          # Display first 10 lines of file.txt
head -n 5 file.txt     # Display first 5 lines
head -c 100 file.txt   # Display first 100 bytes
head file1.txt file2.txt # Display first 10 lines of multiple files
```

5. tail (Display Last Lines of a File)

- **Description:** Displays the last 10 lines of a file by default.
- **Common Options:**
 - `-n N`: Display the last `N` lines.
 - `-c N`: Display the last `N` bytes.
 - `-f`: Follow the file, showing new lines as they are appended (useful for logs).

Examples:

```
tail file.txt          # Display last 10 lines of file.txt
tail -n 5 file.txt     # Display last 5 lines
tail -c 100 file.txt   # Display last 100 bytes
tail -f /var/log/syslog # Follow syslog for real-time updates
```

6. grep (Search Text Using Patterns)

- **Description:** Searches for a pattern in files or input and displays matching lines.
- **Common Options:**
 - `-i`: Ignore case.
 - `-r`: Recursive search in directories.
 - `-n`: Show line numbers.
 - `-v`: Show non-matching lines.
 - `-w`: Match whole words only.
 - `-c`: Count matching lines.

Examples:

```
grep "error" log.txt      # Find lines containing "error" in log.txt
grep -i "error" log.txt   # Case-insensitive search
grep -r "TODO" /home/user # Recursively search for "TODO" in /home/user
grep -n "error" log.txt   # Show line numbers with matches
grep -v "info" log.txt    # Show lines not containing "info"
grep -c "error" log.txt   # Count lines with "error"
```

Practical Tips

- Use `cat` for small files: For quick viewing or concatenation, but avoid for large files.

- **Prefer `less` for large files:** It's more efficient than `cat` and allows navigation.
- **Combine `grep` with pipes:** Use with other commands (e.g., `cat file.txt | grep error`).
- **Monitor logs with `tail -f`:** Useful for real-time log monitoring (e.g., `/var/log/syslog`).
- **Search efficiently:** Use `grep -r` for directories, and combine with `-i` or `-w` for precision.
- **Escape special characters in `grep`:** Use quotes or `\` for patterns with spaces or special characters (e.g., `grep "my pattern" file.txt`).

Practice Tasks

1. Create a file `test.txt` with some text and display its contents using `cat`.
2. View `test.txt` with line numbers using `less`.
3. Display the first 3 lines of `test.txt` using `head`.
4. Display the last 5 lines of `test.txt` using `tail`.
5. Search for the word "test" in `test.txt` with line numbers using `grep`.
6. Combine `tail` and `grep` to show the last 10 lines containing "error" in a log file.

Solution:

```
echo -e "Line 1\ntest line\nLine 3\ntest case\nLine 5" > test.txt
cat test.txt
less -N test.txt          # Press 'q' to quit
head -n 3 test.txt
tail -n 5 test.txt
grep -n "test" test.txt
tail -n 10 /var/log/syslog | grep "error"
```

Linux Process Management (Including PM2)

This document covers Linux commands for managing processes: `ps`, `top`, `htop`, `kill`, `nice`, `fg`, `bg`, and `pm2`. Each command includes a description, common options, and practical examples to help you monitor and control processes effectively, with a focus on both native Linux tools and PM2 for Node.js applications.

1. ps (Process Status)

- **Description:** Displays information about active processes.
- **Common Options:**
 - `-e`: Show all processes.
 - `-f`: Full format listing (includes UID, PID, PPID, etc.).
 - `-u [user]`: Show processes for a specific user.
 - `--forest`: Display process hierarchy.

Examples:

```
ps                # Show processes for the current shell
ps -e             # Show all processes
ps -ef            # Show all processes in full format
ps -u user1       # Show processes for user1
ps -ef --forest   # Show process hierarchy
ps -aux | grep bash # Find bash-related processes
```

2. top (Display System Processes)

- **Description:** Provides a real-time, interactive view of system processes, sorted by resource usage (e.g., CPU, memory).
- **Common Options:**
 - `-u [user]`: Show processes for a specific user.
 - `-d [seconds]`: Set refresh interval.
- **Interactive Commands:**
 - `q`: Quit.
 - `k`: Kill a process (prompts for PID).
 - `r`: Renice a process (prompts for PID and nice value).
 - `1`: Toggle CPU core details.

Examples:

```
top          # Start top with default view
top -u user1  # Show processes for user1
top -d 2      # Refresh every 2 seconds
```

3. htop (Interactive Process Viewer)

- **Description:** A user-friendly, colorful alternative to `top` with mouse/keyboard navigation.
- **Common Options:**
 - `-u [user]`: Show processes for a specific user.
 - `-d [tenths]`: Set refresh interval (in tenths of seconds, e.g., `10` = 1 second).
- **Interactive Commands:**
 - `F2`: Setup (customize display).
 - `F3` or `/`: Search for a process.
 - `F9`: Kill a process.
 - `F7/F8`: Decrease/increase nice value (priority).
 - `q`: Quit.

Examples:

```
htop          # Start htop
htop -u user1  # Show processes for user1
htop -d 20     # Refresh every 2 seconds
```

- **Note:** `htop` may need to be installed (`sudo apt install htop` or equivalent).

4. kill (Terminate Processes)

- **Description:** Sends a signal to terminate or control a process by its PID.
- **Common Signals:**
 - `SIGTERM` (15): Polite termination (default).
 - `SIGKILL` (9): Forceful termination.
 - `SIGHUP` (1): Hangup (reload configuration for some processes).
- **Common Options:**
 - `-l`: List all signal names.
 - `-[signal]`: Specify the signal (e.g., `-9` for `SIGKILL`).

Examples:

```
kill 1234      # Send SIGTERM to process with PID 1234
```

```
kill -9 1234      # Forcefully kill process with PID 1234
kill -HUP 1234    # Send SIGHUP to process with PID 1234
kill -l           # List all available signals
```

5. nice (Set Process Priority)

- **Description:** Sets the priority (nice value) of a new process. Nice values range from -20 (highest priority) to 19 (lowest priority).
- **Common Options:**
 - `-n [value]`: Set the nice value (e.g., `-n 10`).

Examples:

```
nice -n 10 ./script.sh  # Run script.sh with nice value 10
nice -n -5 ./program    # Run program with higher priority (-5)
```

6. renice (Change Running Process Priority)

- **Description:** Changes the priority of an already running process.
- **Common Options:**
 - `[value]`: Set the new nice value.
 - `-p [PID]`: Specify the process ID.
 - `-u [user]`: Change priority for all processes of a user.

Examples:

```
renice 10 -p 1234      # Set nice value 10 for process with PID 1234
renice -5 -p 1234      # Increase priority for PID 1234
renice 0 -u user1       # Reset nice value to 0 for all user1's processes
```

7. fg (Foreground a Process)

- **Description:** brings a background process to the foreground.
- **Common Options:** None typically used, but job number can be specified.

Examples:

```
sleep 100 &          # Start sleep in background
fg                    # Bring most recent background job to foreground
fg %1                 # Bring job number 1 to foreground
```

8. bg (Background a Process)

- **Description:** Resumes a suspended process in the background.
- **Common Options:** None typically used, but job number can be specified.

Examples:

```
sleep 100      # Start sleep, then press Ctrl+Z to suspend
bg             # Resume most recent suspended job in background
bg %1          # Resume job number 1 in background
```

9. pm2 (Process Manager for Node.js Applications)

- **Description:** A production-grade process manager for Node.js (and other languages like Python, Ruby) with built-in load balancing, monitoring, and automatic restarts. It daemonizes applications to keep them running in the background.
- **Common Options/Commands:**
 - `start [app]`: Start an application (e.g., `app.js`).
 - `--name [name]`: Name the process.
 - `-i [n]`: Start in cluster mode with `n` instances (or `max` for CPU core count).
 - `--watch`: Restart on file changes.
 - `list`: List all managed processes.
 - `stop [name|id]`: Stop a process.
 - `restart [name|id]`: Restart a process.
 - `delete [name|id]`: Delete a process from PM2.
 - `logs [name]`: View logs for a process.
 - `monit`: Monitor CPU/memory usage.
 - `startup`: Generate a script to restart PM2 on system boot.
 - `save`: Save current process list for boot persistence.

Examples:

```
npm install pm2@latest -g      # Install PM2 globally
pm2 start app.js --name my-app # Start app.js with name "my-app"
pm2 start app.js -i max        # Start in cluster mode using all CPU cores
pm2 start script.py            # Start a Python script
pm2 list                       # List all managed processes
pm2 stop my-app                # Stop process named my-app
pm2 restart my-app             # Restart process named my-app
pm2 delete my-app              # Delete process named my-app
```

pm2 logs my-app	# View logs for my-app
pm2 monit	# Monitor all processes
pm2 startup	# Generate startup script
pm2 save	# Save process list for boot

- **Note:** PM2 requires Node.js (`sudo apt install nodejs npm` or equivalent). Logs are stored in `~/.pm2/logs/`. Use `pm2 flush` to clear logs.

Practical Tips

- **Find PIDs:** Use `ps -aux | grep [process]` or `pidof [process]` for native processes, or `pm2 list` for PM2-managed processes.
- **Use `htop` for interactivity:** Easier to navigate and kill processes than `top`.
- **Be cautious with `kill -9`:** Use SIGKILL only when SIGTERM fails, as it doesn't allow cleanup.
- **PM2 for production:** Use PM2 for Node.js apps to ensure automatic restarts and load balancing. Use cluster mode (`-i max`) for better performance.
- **Monitor resources:** Use `top`, `htop`, or `pm2 monit` to identify resource-heavy processes.
- **Background tasks:** Use `&` for native background processes, or PM2's daemonization for apps.
- **Check jobs:** Use `jobs` for native background/suspended processes, or `pm2 list` for PM2 processes.

Practice Tasks

1. List all processes for the current user with `ps`.
2. Start `top`, then switch to `htop` to compare interfaces.
3. Start a `sleep 100` process in the background, then bring it to the foreground.
4. Suspend a `sleep 200` process with `Ctrl+Z`, then resume it in the background.
5. Find the PID of a running `bash` process and terminate it with `kill`.
6. Run a script (`script.sh`) with a nice value of 10, then change its priority to 5 using `renice`.
7. Install PM2, start a Node.js app (`app.js`) with name `my-app`, and view its logs.
8. Use PM2 to start a Python script (`script.py`) and monitor it with `pm2 monit`.

Solution:

```
ps -u $USER
```



```
top # Press 'q' to quit
htop # Press 'q' to quit
sleep 100 &
fg
sleep 200 # Press Ctrl+Z to suspend
bg
ps -aux | grep bash # Find PID (e.g., 1234)
kill 1234
nice -n 10 ./script.sh &
renice 5 -p $(pidof script.sh)
npm install pm2@latest -g
pm2 start app.js --name my-app
pm2 logs my-app
pm2 start script.py
pm2 monit
```

Linux Package Management

This document covers Linux package management tools: **apt**, **yum**, **dnf**, **dpkg**, and **rpm**. Each tool includes a description, common options, and practical examples to help you install, remove, and manage software packages on Debian-based (Ubuntu) and Red Hat-based (CentOS, Fedora) systems.

1. apt (Advanced Package Tool)

- **Description:** A high-level package manager for Debian-based systems (e.g., Ubuntu) that handles **.deb** packages and dependencies.
- **Common Options/Commands:**
 - **update:** Update package lists from repositories.
 - **upgrade:** Upgrade installed packages to the latest versions.
 - **install [package]:** Install a package and its dependencies.
 - **remove [package]:** Remove a package, leaving configuration files.
 - **purge [package]:** Remove a package and its configuration files.
 - **autoremove:** Remove unneeded dependencies.
 - **search [term]:** Search for packages by name or description.

Examples:

```
sudo apt update          # Refresh package lists
sudo apt upgrade         # Upgrade all installed packages
sudo apt install vim     # Install vim
```

```
sudo apt remove vim          # Remove vim, keep config files
sudo apt purge vim           # Remove vim and its config files
sudo apt autoremove          # Remove unneeded dependencies
sudo apt search nginx        # Search for nginx-related packages
```

- **Note:** Requires `sudo` for administrative tasks. Use `apt-cache` for additional queries (e.g., `apt-cache show [package]`).

2. yum (Yellowdog Updater, Modified)

- **Description:** A package manager for older Red Hat-based systems (e.g., CentOS 7) that handles `.rpm` packages and dependencies.
- **Common Options/Commands:**
 - `update`: Update package lists and upgrade packages.
 - `install [package]`: Install a package and its dependencies.
 - `remove [package]`: Remove a package.
 - `search [term]`: Search for packages by name or description.
 - `info [package]`: Show package details.
 - `clean all`: Clear cached data.

Examples:

```
sudo yum update              # Update package lists and upgrade
sudo yum install httpd       # Install Apache web server
sudo yum remove httpd        # Remove Apache
sudo yum search vim           # Search for vim-related packages
sudo yum info httpd          # Show details about httpd
sudo yum clean all           # Clear cached data
```

- **Note:** Requires `sudo`. Largely replaced by `dnf` in newer Red Hat-based systems.

3. dnf (Dandified YUM)

- **Description:** A modern replacement for `yum` on Red Hat-based systems (e.g., Fedora, CentOS 8+), handling `.rpm` packages with improved performance.
- **Common Options/Commands:**
 - `update` or `upgrade`: Update package lists and upgrade packages.
 - `install [package]`: Install a package and its dependencies.
 - `remove [package]`: Remove a package.

- `search [term]`: Search for packages.
- `info [package]`: Show package details.
- `autoremove`: Remove unneeded dependencies.
- `clean all`: Clear cached data.

Examples:

```
sudo dnf update           # Update package lists and upgrade
sudo dnf install nginx    # Install nginx
sudo dnf remove nginx     # Remove nginx
sudo dnf search vim       # Search for vim-related packages
sudo dnf info nginx       # Show details about nginx
sudo dnf autoremove       # Remove unneeded dependencies
sudo dnf clean all        # Clear cached data
```

- **Note:** Requires `sudo`. Similar syntax to `yum` but with better dependency resolution.

4. dpkg (Debian Package)

- **Description:** A low-level tool for handling `.deb` packages on Debian-based systems. Used directly for installing or querying individual packages.
- **Common Options:**
 - `-i [package.deb]`: Install a `.deb` package.
 - `-r [package]`: Remove a package, leaving configuration files.
 - `-P [package]`: Purge a package and its configuration files.
 - `-l [pattern]`: List installed packages matching a pattern.
 - `-s [package]`: Show package status/details.

Examples:

```
sudo dpkg -i package.deb    # Install a .deb package
sudo dpkg -r vim            # Remove vim, keep config files
sudo dpkg -P vim            # Purge vim and its config files
dpkg -l | grep vim          # List installed vim packages
dpkg -s vim                 # Show status of vim package
```

- **Note:** Requires `sudo` for installation/removal. Use `apt` for dependency management, as `dpkg` doesn't handle dependencies automatically.

5. rpm (RPM Package Manager)

- **Description:** A low-level tool for handling `.rpm` packages on Red Hat-based systems. Used for installing or querying individual packages.
- **Common Options:**
 - `-i [package.rpm]`: Install a `.rpm` package.
 - `-e [package]`: Remove a package.
 - `-U [package.rpm]`: Upgrade or install a package.
 - `-qa [pattern]`: Query all installed packages matching a pattern.
 - `-qi [package]`: Show package information.

Examples:

```
sudo rpm -i package.rpm      # Install a .rpm package
sudo rpm -e httpd            # Remove httpd package
sudo rpm -U package.rpm      # Upgrade or install package.rpm
rpm -qa | grep httpd         # List installed httpd packages
rpm -qi httpd                # Show details about httpd
```

- **Note:** Requires `sudo` for installation/removal. Use `yum` or `dnf` for dependency management, as `rpm` doesn't handle dependencies automatically.

Practical Tips

- **Update regularly:** Run `apt update && apt upgrade`, `yum update`, or `dnf update` to keep systems secure.
- **Use high-level tools:** Prefer `apt`, `yum`, or `dnf` over `dpkg` or `rpm` for easier dependency management.
- **Search before installing:** Use `apt search`, `yum search`, or `dnf search` to find package names.
- **Clean up:** Use `autoremove` and `clean` commands to free disk space.
- **Handle dependency issues:** If `dpkg` or `rpm` fails due to dependencies, use `apt` or `dnf` to resolve them.
- **Check package status:** Use `dpkg -s` or `rpm -qi` to verify installed packages.

Practice Tasks

1. Update package lists and upgrade all packages on a Debian-based system using `apt`.
2. Install `nginx` on a Red Hat-based system using `dnf`.
3. Search for `vim` packages using `yum` or `dnf`.
4. Install a `.deb` package (`package.deb`) using `dpkg`.
5. Remove a package (`httpd`) using `rpm`.

6. List all installed packages containing "nginx" using `dpkg` or `rpm`.

Solution:

```
sudo apt update && sudo apt upgrade
sudo dnf install nginx
sudo yum search vim # Or sudo dnf search vim
sudo dpkg -i package.deb
sudo rpm -e httpd
dpkg -l | grep nginx # Or rpm -qa | grep nginx
```

Linux Network Utilities

This document covers Linux commands for network diagnostics and management: `ping`, `netstat`, `ss`, `ip`, `nmcli`, `curl`, and `wget`. Each command includes a description, common options, and practical examples to help you troubleshoot and interact with networks effectively. A comparison between `curl` and `wget` is also included.

1. ping (Packet Internet Groper)

- **Description:** Tests network connectivity by sending ICMP echo requests to a host and measuring response time.
- **Common Options:**
 - `-c [count]`: Send a specified number of packets.
 - `-i [interval]`: Set interval between packets (in seconds).
 - `-s [size]`: Set packet size (in bytes).
 - `-t [ttl]`: Set time-to-live for packets.

Examples:

```
ping google.com          # Ping google.com continuously
ping -c 4 google.com     # Send 4 ping requests
ping -i 0.5 google.com   # Ping every 0.5 seconds
ping -s 100 192.168.1.1  # Ping with 100-byte packets
ping -t 64 8.8.8.8       # Ping with TTL of 64
```

- **Note:** Use `Ctrl+C` to stop continuous pinging.

2. netstat (Network Statistics)

- **Description:** Displays network connections, routing tables, and interface statistics. Often replaced by `ss` in modern systems.
- **Common Options:**
 - `-a`: Show all connections (including listening).
 - `-t`: Show TCP connections.
 - `-u`: Show UDP connections.
 - `-l`: Show listening sockets.
 - `-p`: Show program/PID associated with connections.
 - `-n`: Show numerical addresses (not resolved to hostnames).

Examples:

```
netstat -tuln      # Show listening TCP/UDP ports
netstat -anp      # Show all connections with programs/PIDs
netstat -r        # Show routing table
netstat -i        # Show interface statistics
```

-
- **Note:** May require installation (`sudo apt install net-tools`).

3. ss (Socket Statistics)

- **Description:** A modern replacement for `netstat`, providing detailed information about network sockets.
- **Common Options:**
 - `-t`: Show TCP sockets.
 - `-u`: Show UDP sockets.
 - `-l`: Show listening sockets.
 - `-a`: Show all sockets (listening and non-listening).
 - `-p`: Show process/PID using the socket.
 - `-n`: Show numerical addresses.

Examples:

```
ss -tuln      # Show listening TCP/UDP ports
ss -anp      # Show all sockets with programs/PIDs
ss -t -a      # Show all TCP sockets
ss -s        # Show socket summary statistics
```

4. ip (IP Configuration)

- **Description:** Manages network interfaces, routes, and IP addresses. Part of the `iproute2` suite, replacing `ifconfig`.
- **Common Commands:**
 - `ip addr`: Show IP addresses and interfaces.
 - `ip link`: Show network interfaces and their status.
 - `ip route`: Show or manage routing table.
- **Common Options:**
 - `show`: Display information (default).
 - `add/del`: Add or delete configurations.

Examples:

```
ip addr show          # Show all network interfaces and IPs
ip link show          # Show interface status
ip route show         # Show routing table
sudo ip addr add 192.168.1.100/24 dev eth0 # Add IP to interface eth0
sudo ip link set eth0 up # Bring interface eth0 up
```

5. nmcli (NetworkManager Command Line Interface)

- **Description:** Manages network connections on systems using NetworkManager (common in Ubuntu, Fedora).
- **Common Commands:**
 - `nmcli connection show`: List network connections.
 - `nmcli device status`: Show device status.
 - `nmcli con up [name]`: Activate a connection.
 - `nmcli con down [name]`: Deactivate a connection.

Examples:

```
nmcli connection show # List all network connections
nmcli device status    # Show status of network devices
nmcli con up my-wifi    # Activate connection named my-wifi
nmcli con down my-wifi  # Deactivate connection named my-wifi
nmcli con add type wifi con-name my-wifi ssid MySSID # Add a Wi-Fi connection
```

- **Note:** Requires NetworkManager (`sudo apt install network-manager`).

6. curl (Client URL)

- **Description:** Transfers data to/from a server, supporting protocols like HTTP, HTTPS, FTP, and more. Ideal for scripting and API interactions.
- **Common Options:**
 - `-o [file]`: Save output to a file.
 - `-I`: Fetch headers only.
 - `-X [method]`: Specify HTTP method (e.g., GET, POST).
 - `-d [data]`: Send data in a POST request.
 - `-L`: Follow redirects.

Examples:

```
curl http://example.com      # Fetch content from example.com
curl -o output.html http://example.com # Save content to output.html
curl -I http://example.com  # Fetch headers only
curl -X POST -d "name=value" http://example.com/api # Send POST request
curl -L http://example.com  # Follow redirects
```

7. wget (Web Get)

- **Description:** Downloads files from the web, supporting HTTP, HTTPS, and FTP. Designed for robust file downloads, including recursive fetching.
- **Common Options:**
 - `-O [file]`: Save output to a specific file.
 - `-c`: Resume a partially downloaded file.
 - `-r`: Recursive download (e.g., entire website).
 - `-q`: Quiet mode (no output).
 - `--limit-rate=[rate]`: Limit download speed (e.g., 100k).

Examples:

```
wget http://example.com/file.txt # Download file.txt
wget -O myfile.txt http://example.com/file.txt # Save as myfile.txt
wget -c http://example.com/largefile.zip # Resume partial download
wget -r http://example.com          # Recursively download website
wget --limit-rate=100k http://example.com/file.zip # Limit speed to 100 KB/s
```


Comparison: curl vs. wget

- **Purpose:**
 - **curl**: Designed for transferring data, including API requests, with support for multiple protocols (HTTP, HTTPS, FTP, SFTP, etc.). Best for scripting and complex requests.
 - **wget**: Focused on downloading files and recursive website fetching. Ideal for simple downloads and mirroring websites.
- **Protocols:**
 - **curl**: Supports a broader range of protocols (e.g., HTTP, HTTPS, FTP, SFTP, LDAP, IMAP).
 - **wget**: Limited to HTTP, HTTPS, and FTP.
- **Output Handling:**
 - **curl**: Outputs to **stdout** by default, requiring **-o** or **>** for file saving.
 - **wget**: Saves to a file by default, with automatic filename detection.
- **Interactivity:**
 - **curl**: Better for scripting, API calls, and non-interactive tasks (e.g., POST requests, authentication).
 - **wget**: Better for recursive downloads and resuming interrupted downloads.
- **Features:**
 - **curl**: Supports advanced HTTP features (e.g., POST, PUT, headers, cookies, authentication).
 - **wget**: Supports recursive downloads (**-r**) and bandwidth limiting (**--limit-rate**).
- **Examples:**

Use **curl** for API testing:

```
curl -X POST -H "Content-Type: application/json" -d '{"key":"value"}' http://api.example.com
```

Use **wget** for recursive website download:

```
wget -r http://example.com
```

- **When to Use:**
 - Choose **curl** for API interactions, custom HTTP requests, or non-HTTP protocols.
 - Choose **wget** for simple file downloads, resuming large downloads, or mirroring websites.

Practical Tips

- **Use `ping` for quick checks:** Test connectivity before deeper troubleshooting.
- **Prefer `ss` over `netstat`:** `ss` is faster and more modern, but `netstat` may be used on older systems.
- **Combine `ip` with `nmcli`:** Use `ip` for low-level interface details and `nmcli` for managing NetworkManager connections.
- **Secure `curl` and `wget`:** Use HTTPS URLs and verify certificates (e.g., `curl --cacert [file]` or `wget --no-check-certificate` for testing).
- **Pipe `curl` output:** Combine with `grep` or `jq` for parsing (e.g., `curl api.example.com | jq .`).
- **Monitor with `watch`:** Use `watch ss -tuln` to monitor network sockets in real-time.

Practice Tasks

1. Ping `google.com` 4 times and check the response time.
2. List all listening TCP/UDP ports using `ss`.
3. Display the routing table using `ip`.
4. Show all network connections with `nmcli`.
5. Download a file from `http://example.com/file.txt` using `wget`.
6. Fetch the headers of `http://example.com` using `curl`.
7. Use `curl` to send a POST request with data `name=test` to `http://example.com/api`.
8. Use `wget` to resume a partial download of a large file.

Solution:

```
ping -c 4 google.com
ss -tuln
ip route show
nmcli connection show
wget http://example.com/file.txt
curl -I http://example.com
curl -X POST -d "name=test" http://example.com/api
wget -c http://example.com/largefile.zip
```

Linux DNS and Hosts File

This document covers the configuration and management of DNS-related files and concepts in Linux: `/etc/hosts`, `/etc/resolv.conf`, and DNS configuration basics. Each section includes a description, common configurations, and practical examples to help you understand and manage DNS settings effectively.

1. `/etc/hosts` (Local Hosts File)

- **Description:** Maps hostnames to IP addresses locally, overriding DNS lookups. Used for static hostname resolution without querying a DNS server.
- **File Location:** `/etc/hosts`

Format:

IP_address hostname [aliases...]

-
- **Common Entries:**
 - `127.0.0.1 localhost`: Maps localhost to the loopback address.
 - `::1 localhost`: IPv6 loopback address.
 - Custom mappings for local networks or testing (e.g., `192.168.1.100 myserver`).

Examples:

View contents of `/etc/hosts`

```
cat /etc/hosts
```

Example output:

```
# 127.0.0.1 localhost
```

```
# ::1 localhost
```

```
# 192.168.1.100 myserver.local myserver
```

Add a new entry (requires sudo)

```
sudo nano /etc/hosts
```

Add: `192.168.1.101 webserver.local webserver`

Save and exit

Test resolution

```
ping webserver.local # Should resolve to 192.168.1.101
```

- **Note:** Requires `sudo` to edit. Used for local overrides or when DNS is unavailable.

2. /etc/resolv.conf (Resolver Configuration File)

- **Description:** Specifies DNS servers and domains for name resolution. Managed automatically by tools like NetworkManager in modern systems.
- **File Location:** `/etc/resolv.conf`
- **Common Directives:**
 - `nameserver [IP]`: Specifies a DNS server (e.g., `8.8.8.8` for Google DNS).
 - `search [domain]`: Appends domain suffixes for short hostnames (e.g., `search example.com`).
 - `domain [domain]`: Sets the local domain name.

Examples:

```
# View contents of /etc/resolv.conf
```

```
cat /etc/resolv.conf
```

```
# Example output:
```

```
# nameserver 8.8.8.8
```

```
# nameserver 8.8.4.4
```

```
# search mydomain.local
```

```
# Manually add a DNS server (requires sudo)
```

```
sudo nano /etc/resolv.conf
```

```
# Add: nameserver 1.1.1.1
```

```
# Save and exit
```

```
# Test DNS resolution
```

```
nslookup google.com # Should use specified DNS servers
```

- **Note:** Changes to `/etc/resolv.conf` may be overwritten by NetworkManager or DHCP. To persist changes, configure the network manager (e.g., `nmccli`) or edit `/etc/resolvconf/resolv.conf.d/base` on some systems.

3. DNS Configuration Basics

- **Description:** DNS (Domain Name System) resolves hostnames to IP addresses using external servers. Linux systems rely on `/etc/resolv.conf` or tools like `systemd-resolved` for DNS queries.
- **Key Concepts:**
 - **Nameservers:** IP addresses of DNS servers (e.g., `8.8.8.8` for Google, `1.1.1.1` for Cloudflare).
 - **DNS Queries:** Tools like `nslookup`, `dig`, or `host` test resolution.

- **Caching:** Systems like `systemd-resolved` cache DNS responses to improve performance.
- **Tools:**
 - `nslookup`: Query DNS servers interactively.
 - `dig`: Detailed DNS query tool.
 - `host`: Simple DNS lookup tool.

Examples:

Test DNS resolution with nslookup

nslookup google.com

Example output: Resolves to IPs like 142.250.190.78

Use dig for detailed DNS query

dig google.com

Example output: Shows query time, server used, and DNS records

Simple lookup with host

host google.com

Example output: google.com has address 142.250.190.78

Check systemd-resolved status (if applicable)

systemctl status systemd-resolved

Enable DNS caching

sudo systemctl enable systemd-resolved

sudo systemctl start systemd-resolved

-
- **Note:** Install `dnsutils` for `nslookup`, `dig`, and `host` (`sudo apt install dnsutils` or equivalent).

Practical Tips

- **Use `/etc/hosts` for testing:** Add entries to simulate DNS for development or debugging (e.g., mapping a domain to a local IP).
- **Avoid manual `/etc/resolv.conf` edits:** Use network management tools like `nmcli` to persist DNS settings.
- **Test DNS:** Use `nslookup` or `dig` to verify DNS server functionality.
- **Fallback DNS:** Add reliable public DNS servers (e.g., `8.8.8.8`, `1.1.1.1`) if local DNS fails.
- **Check for `systemd-resolved`:** Modern systems use it for DNS caching; check `/etc/systemd/resolved.conf` for advanced settings.

- **Backup before editing:** Copy `/etc/hosts` or `/etc/resolv.conf` before changes (`sudo cp /etc/hosts /etc/hosts.bak`).

Practice Tasks

1. View the contents of `/etc/hosts` and add a mapping for `192.168.1.200` to `testserver.local`.
2. Check the DNS servers in `/etc/resolv.conf`.
3. Add `1.1.1.1` as a DNS server to `/etc/resolv.conf` (if not managed by NetworkManager).
4. Use `nslookup` to resolve `example.com`.
5. Use `dig` to query DNS details for `google.com`.
6. Check if `systemd-resolved` is running and enable it if not.

Solution:

```
cat /etc/érêt
```

System: You are Grok 3 built by xAI.

The current date and time is 04:52 PM AEST on Friday, June 27, 2025.

Continuing from where the previous response left off, the Markdown document for the **DNS and Hosts File** topic was being created but was cut off mid-example. Below is the completed Markdown document covering `/etc/hosts`, `/etc/resolv.conf`, and DNS configuration basics, with descriptions, common configurations, practical examples, and the completed practice task solution. You can copy and paste this directly into a Markdown file (e.g., `dns_and_hosts.md`) for your study notes.

```
<xaiArtifact artifact_id="23698488-8e92-4ca0-84d8-6a1fa10ca8f9"
artifact_version_id="a0d8827a-2fb7-4e25-83c1-933c2b15bcb6" title="dns_and_hosts.md"
contentType="text/markdown">
```

Linux DNS and Hosts File

This document covers the configuration and management of DNS-related files and concepts in Linux: `/etc/hosts`, `/etc/resolv.conf`, and DNS configuration basics. Each section includes a description, common configurations, and practical examples to help you understand and manage DNS settings effectively.

1. /etc/hosts (Local Hosts File)

- **Description**: Maps hostnames to IP addresses locally, overriding DNS lookups. Used for static hostname resolution without querying a DNS server.

- **File Location**: `/etc/hosts`
- **Format**:

IP_address hostname [aliases...]

- **Common Entries**:
- `127.0.0.1 localhost`: Maps localhost to the loopback address.
- `::1 localhost`: IPv6 loopback address.
- Custom mappings for local networks or testing (e.g., `192.168.1.100 myserver`).
- **Examples**:

```
``bash
# View contents of /etc/hosts
cat /etc/hosts
# Example output:
# 127.0.0.1 localhost
# ::1 localhost
# 192.168.1.100 myserver.local myserver

# Add a new entry (requires sudo)
sudo nano /etc/hosts
# Add: 192.168.1.101 webserver.local webserver
# Save and exit

# Test resolution
ping webserver.local # Should resolve to 192.168.1.101
```

- **Note**: Requires `sudo` to edit. Used for local overrides or when DNS is unavailable.

2. `/etc/resolv.conf` (Resolver Configuration File)

- **Description**: Specifies DNS servers and domains for name resolution. Managed automatically by tools like NetworkManager in modern systems.
- **File Location**: `/etc/resolv.conf`
- **Common Directives**:
 - `nameserver [IP]`: Specifies a DNS server (e.g., `8.8.8.8` for Google DNS).
 - `search [domain]`: Appends domain suffixes for short hostnames (e.g., `search example.com`).
 - `domain [domain]`: Sets the local domain name.

Examples:

```
# View contents of /etc/resolv.conf
```

```
cat /etc/resolv.conf
# Example output:
# nameserver 8.8.8.8
# nameserver 8.8.4.4
# search mydomain.local

# Manually add a DNS server (requires sudo)
sudo nano /etc/resolv.conf
# Add: nameserver 1.1.1.1
# Save and exit

# Test DNS resolution
nslookup google.com # Should use specified DNS servers
```

-
- **Note:** Changes to `/etc/resolv.conf` may be overwritten by NetworkManager or DHCP. To persist changes, configure the network manager (e.g., `nmcli`) or edit `/etc/resolvconf/resolv.conf.d/base` on some systems.

3. DNS Configuration Basics

- **Description:** DNS (Domain Name System) resolves hostnames to IP addresses using external servers. Linux systems rely on `/etc/resolv.conf` or tools like `systemd-resolved` for DNS queries.
- **Key Concepts:**
 - **Nameservers:** IP addresses of DNS servers (e.g., `8.8.8.8` for Google, `1.1.1.1` for Cloudflare).
 - **DNS Queries:** Tools like `nslookup`, `dig`, or `host` test resolution.
 - **Caching:** Systems like `systemd-resolved` cache DNS responses to improve performance.
- **Tools:**
 - `nslookup`: Query DNS servers interactively.
 - `dig`: Detailed DNS query tool.
 - `host`: Simple DNS lookup tool.

Examples:

```
# Test DNS resolution with nslookup
nslookup google.com
# Example output: Resolves to IPs like 142.250.190.78

# Use dig for detailed DNS query
dig google.com
```


Example output: Shows query time, server used, and DNS records

Simple lookup with host

host google.com

Example output: google.com has address 142.250.190.78

Check systemd-resolved status (if applicable)

systemctl status systemd-resolved

Enable DNS caching

sudo systemctl enable systemd-resolved

sudo systemctl start systemd-resolved

-
- **Note:** Install `dnsutils` for `nslookup`, `dig`, and `host` (`sudo apt install dnsutils` or equivalent).

Practical Tips

- **Use `/etc/hosts` for testing:** Add entries to simulate DNS for development or debugging (e.g., mapping a domain to a local IP).
- **Avoid manual `/etc/resolv.conf` edits:** Use network management tools like `nmcli` to persist DNS settings.
- **Test DNS:** Use `nslookup` or `dig` to verify DNS server functionality.
- **Fallback DNS:** Add reliable public DNS servers (e.g., `8.8.8.8`, `1.1.1.1`) if local DNS fails.
- **Check for `systemd-resolved`:** Modern systems use it for DNS caching; check `/etc/systemd/resolved.conf` for advanced settings.
- **Backup before editing:** Copy `/etc/hosts` or `/etc/resolv.conf` before changes (`sudo cp /etc/hosts /etc/hosts.bak`).

Practice Tasks

1. View the contents of `/etc/hosts` and add a mapping for `192.168.1.200` to `testserver.local`.
2. Check the DNS servers in `/etc/resolv.conf`.
3. Add `1.1.1.1` as a DNS server to `/etc/resolv.conf` (if not managed by NetworkManager).
4. Use `nslookup` to resolve `example.com`.
5. Use `dig` to query DNS details for `google.com`.
6. Check if `systemd-resolved` is running and enable it if not.

Solution:

```
cat /etc/hosts
sudo nano /etc/hosts
# Add: 192.168.1.200 testserver.local testserver
# Save and exit
cat /etc/resolv.conf
sudo nano /etc/resolv.conf
# Add: nameserver 1.1.1.1
# Save and exit
nslookup example.com
dig google.com
systemctl status systemd-resolved
sudo systemctl enable systemd-resolved
sudo systemctl start systemd-resolved
```

Linux Cron Jobs and Scheduling

This document covers Linux tools for scheduling tasks: `crontab`, `at`, and automating tasks. Each section includes a description, common syntax, and practical examples to help you automate repetitive tasks effectively.

1. crontab (Cron Table)

- **Description:** Schedules recurring tasks using cron, a time-based job scheduler. Each user has their own crontab file.
- **File Location:** User crontabs are stored in `/var/spool/cron/crontabs/`, system-wide in `/etc/crontab` or `/etc/cron.*`.

Syntax:

```
* * * * * /path/to/script.sh
```

Fields: minute (0-59), hour (0-23), day (1-31), month (1-12), day-of-week (0-7)

-
- **Common Commands:**
 - `crontab -e`: Edit the current user's crontab.
 - `crontab -l`: List the current user's crontab.
 - `crontab -r`: Remove the current user's crontab.
- **Special Strings:**
 - `@reboot`: Run once at system startup.

- `@daily`, `@weekly`, `@monthly`: Run daily, weekly, or monthly.

Examples:

Edit crontab

`crontab -e`

Add: Run backup.sh every day at 2 AM

`0 2 * * * /home/user/backup.sh`

Run script every 5 minutes

`* /5 * * * * /home/user/check_status.sh`

Run script at system reboot

`@reboot /home/user/start_app.sh`

List crontab

`crontab -l`

Example output:

`# 0 2 * * * /home/user/backup.sh`

`# * /5 * * * * /home/user/check_status.sh`

Remove crontab

`crontab -r`

-
- **Note:** Ensure scripts have executable permissions (`chmod +x script.sh`). Logs are typically in `/var/log/syslog` or `/var/log/cron`.

2. at (One-Time Task Scheduling)

- **Description:** Schedules a one-time task to run at a specified time.
- **Common Commands:**
 - `at [time]`: Schedule a task for a specific time.
 - `atq`: List pending `at` jobs.
 - `atrm [job_id]`: Remove a scheduled `at` job.
- **Time Formats:**
 - Absolute: `HH:MM` (e.g., `14:30`), `YYYY-MM-DD HH:MM` (e.g., `2025-06-28 14:30`).
 - Relative: `now + [time]` (e.g., `now + 10 minutes`).

Examples:

Schedule a script to run at 3 PM today

at 15:00

Prompt appears, enter command:

```
/home/user/script.sh
# Press Ctrl+D to save

# Schedule a command for 10 minutes from now
echo "/home/user/backup.sh" | at now + 10 minutes

# List pending at jobs
atq
# Example output: 1 Fri Jun 27 15:00:00 2025 a user

# Remove job with ID 1
atrm 1
```

-
- **Note:** Requires the `at` package (`sudo apt install at`). Jobs are executed in the user's shell environment.

3. Automating Tasks

- **Description:** Combine `crontab` and `at` with scripts to automate repetitive or one-time tasks, such as backups, system monitoring, or cleanup.
- **Key Practices:**
 - Write scripts with absolute paths (e.g., `/usr/bin/echo` instead of `echo`).
 - Redirect output to logs for debugging (e.g., `>> /var/log/myscript.log 2>&1`).
 - Test scripts manually before scheduling.
 - Use environment variables in crontab if needed (e.g., `PATH=/usr/bin:/bin`).

Examples:

```
# Create a backup script
nano backup.sh
# Content:
#!/bin/bash
tar -czf /home/user/backup_$(date +%F).tar.gz /home/user/docs >> /var/log/backup.log 2>&1

# Make executable
chmod +x backup.sh

# Schedule daily at 1 AM via crontab
crontab -e
# Add: 0 1 * * * /home/user/backup.sh

# Schedule a one-time cleanup 2 hours from now
```

```
echo "rm -f /tmp/*.log" | at now + 2 hours
```

-

Practical Tips

- **Test scripts first:** Run scripts manually to ensure they work before scheduling.
- **Log output:** Redirect script output to a log file for troubleshooting (e.g., `>> /logfile 2>&1`).
- **Use absolute paths:** Avoid errors in cron due to limited environment variables.
- **Check cron logs:** Look in `/var/log/syslog` or `/var/log/cron` for cron job execution details.
- **Limit `at` usage:** Use `at` for one-time tasks and `crontab` for recurring tasks.
- **Secure scripts:** Ensure scripts have proper permissions and avoid sensitive data in plaintext.
- **Cron syntax tools:** Use online tools like `crontab.guru` to verify cron schedules.

Practice Tasks

1. Create a script `log_date.sh` that logs the current date to `/tmp/date.log`.
2. Schedule `log_date.sh` to run every minute using `crontab`.
3. List the current user's crontab.
4. Schedule a one-time task to run `echo "Hello" > /tmp/hello.txt` 5 minutes from now using `at`.
5. List all pending `at` jobs and remove one if it exists.
6. Verify the output of the cron job in `/tmp/date.log`.

Solution:

```
# Task 1: Create script
cat << 'EOF' > log_date.sh
#!/bin/bash
echo "Date: $(date)" >> /tmp/date.log
EOF
chmod +x log_date.sh
```

```
# Task 2: Schedule in crontab
crontab -e
# Add: * * * * * /home/user/log_date.sh
# Save and exit
```

```
# Task 3: List crontab
```

```
crontab -l
```

```
# Task 4: Schedule one-time task
```

```
echo "echo 'Hello' > /tmp/hello.txt" | at now + 5 minutes
```

```
# Task 5: List and remove at jobs
```

```
atq
```

```
# If job ID is 1:
```

```
atrm 1
```

```
# Task 6: Verify cron output
```

```
cat /tmp/date.log
```

Linux Systemd and Services

This document covers tools for managing system services in Linux using **systemd**: **systemctl**, **service**, and **journalctl**. Each section includes a description, common commands, and practical examples to help you manage and troubleshoot system services effectively.

1. systemctl (System Control)

- **Description:** Manages **systemd** services, units, and system state. **systemd** is the init system for most modern Linux distributions (e.g., Ubuntu, Fedora).
- **Common Commands:**
 - **start [unit]**: Start a service.
 - **stop [unit]**: Stop a service.
 - **restart [unit]**: Restart a service.
 - **enable [unit]**: Enable a service to start at boot.
 - **disable [unit]**: Disable a service from starting at boot.
 - **status [unit]**: Show service status.
 - **list-units**: List all loaded units.
 - **list-unit-files**: List all unit files and their state.

Examples:

```
# Start a service
```

```
sudo systemctl start nginx
```

```
# Stop a service
```

```
sudo systemctl stop nginx
```

```
# Restart a service
sudo systemctl restart nginx
```

```
# Enable service to start at boot
sudo systemctl enable nginx
```

```
# Disable service from starting at boot
sudo systemctl disable nginx
```

```
# Check service status
systemctl status nginx
# Example output: Shows active/inactive state, PID, and logs
```

```
# List all active units
systemctl list-units --type=service
```

```
# List all unit files
systemctl list-unit-files --type=service
```

-
- **Note:** Requires `sudo` for administrative tasks. Unit types include `service`, `timer`, `socket`, etc.

2. service (Service Management)

- **Description:** A legacy wrapper script for managing services, compatible with both `systemd` and older init systems (e.g., SysVinit). Often redirects to `systemctl` on `systemd`-based systems.
- **Common Commands:**
 - `start [service]`: Start a service.
 - `stop [service]`: Stop a service.
 - `restart [service]`: Restart a service.
 - `status [service]`: Show service status.

Examples:

```
# Start a service
sudo service ssh start
```

```
# Stop a service
sudo service ssh stop
```

```
# Restart a service
```

```
sudo service ssh restart
```

```
# Check service status
```

```
service ssh status
```

```
# Example output: Shows running/stopped state and recent logs
```

-
- **Note:** Requires `sudo` for administrative tasks. Use `systemctl` for more advanced `systemd` features.

3. journalctl (Systemd Journal Logs)

- **Description:** Queries and displays logs from the `systemd` journal, which collects logs from services and the system.
- **Common Options:**
 - `-u [unit]`: Show logs for a specific service.
 - `-b`: Show logs from the current boot.
 - `-f`: Follow logs in real-time (like `tail -f`).
 - `-n [lines]`: Show the last `n` lines.
 - `--since [time]`: Show logs since a specific time (e.g., `2025-06-27`).
 - `-r`: Reverse order (newest first).

Examples:

```
# View all journal logs
```

```
journalctl
```

```
# View logs for a specific service
```

```
journalctl -u nginx
```

```
# View logs from current boot
```

```
journalctl -b
```

```
# Follow logs in real-time
```

```
journalctl -f
```

```
# Show last 10 lines of logs
```

```
journalctl -n 10
```

```
# Show logs since a specific date
```

```
journalctl --since "2025-06-27"
```

```
# Show logs in reverse order
```

```
journalctl -r
```


-
- **Note:** Requires `sudo` to view all logs. Logs are stored in `/var/log/journal/` (if persistent journaling is enabled).

Practical Tips

- **Prefer `systemctl` over `service`:** `systemctl` provides more control and is native to `systemd`.
- **Check service status:** Use `systemctl status` or `service status` to troubleshoot service issues.
- **Monitor logs in real-time:** Use `journalctl -f` for debugging running services.
- **Persistent journaling:** Enable with `sudo mkdir -p /var/log/journal` and `sudo systemd-tmpfiles --create --prefix /var/log/journal`.
- **Filter logs:** Combine `journalctl` options (e.g., `journalctl -u nginx -b` for nginx logs from current boot).
- **Restart safely:** Use `restart` instead of `stop` then `start` to minimize downtime.

Practice Tasks

1. Start the `ssh` service using `systemctl`.
2. Check the status of the `nginx` service using `service`.
3. Enable the `nginx` service to start at boot.
4. View the last 20 lines of logs for the `ssh` service using `journalctl`.
5. Monitor real-time logs for the `nginx` service.
6. List all active services using `systemctl`.

Solution:

```
sudo systemctl start ssh
service nginx status
sudo systemctl enable nginx
journalctl -u ssh -n 20
journalctl -u nginx -f
systemctl list-units --type=service
```

Linux User Management

This document covers Linux commands for managing users: `adduser`, `usermod`, `passwd`, and `groups`. Each command includes a description, common options, and practical examples to help you create, modify, and manage user accounts effectively.

1. adduser (Add a User)

- **Description:** An interactive, high-level command to create a new user, setting up their home directory, password, and other details.
- **Common Options:**
 - `--home [dir]`: Specify the home directory.
 - `--shell [shell]`: Set the user's default shell (e.g., `/bin/bash`).
 - `--ingroup [group]`: Set the primary group.
 - `--gecos [info]`: Set user information (e.g., full name).

Examples:

Add a new user interactively

```
sudo adduser newuser
```

Follow prompts to set password, full name, etc.

Add a user with specific options

```
sudo adduser --home /home/newuser --shell /bin/zsh --ingroup developers newuser
```

Verify user creation

```
id newuser
```

Example output: uid=1001(newuser) gid=1001(developers) groups=1001(developers)

-

- **Note:** Requires `sudo`. Creates home directory and adds user to `/etc/passwd` and `/etc/shadow`. Use `useradd` for a lower-level alternative.

2. usermod (Modify a User)

- **Description:** Modifies an existing user's account settings, such as group membership, home directory, or shell.
- **Common Options:**
 - `-aG [group]`: Add user to supplementary groups (use with `-a` to append).
 - `-d [dir]`: Change home directory.
 - `-s [shell]`: Change login shell.
 - `-l [newname]`: Change username.

- **-L**: Lock the user account (disable login).
- **-U**: Unlock the user account.

Examples:

Add user to a group

```
sudo usermod -aG sudo newuser
```

Verify group membership

```
groups newuser
```

Output: newuser : developers sudo

Change home directory

```
sudo usermod -d /newhome/newuser newuser
```

Change shell

```
sudo usermod -s /bin/bash newuser
```

Rename user

```
sudo usermod -l newname newuser
```

Lock user account

```
sudo usermod -L newuser
```

Unlock user account

```
sudo usermod -U newuser
```

-
- **Note:** Requires **sudo**. Use **-a** with **-G** to avoid overwriting existing groups.

3. passwd (Change User Password)

- **Description:** Changes a user's password or manages password settings.
- **Common Options:**
 - **[username]**: Specify the user (default is current user).
 - **-l**: Lock the password (disable login).
 - **-u**: Unlock the password.
 - **-d**: Delete the password (allow passwordless login).

Examples:

Change current user's password

```
passwd
```

Follow prompts to enter new password

Change another user's password

```
sudo passwd newuser
```

```
# Follow prompts
```

```
# Lock a user's password
```

```
sudo passwd -l newuser
```

```
# Unlock a user's password
```

```
sudo passwd -u newuser
```

```
# Delete a user's password
```

```
sudo passwd -d newuser
```

-
- **Note:** Requires `sudo` for other users. Locked accounts show `!` in `/etc/shadow`.

4. groups (Display User Groups)

- **Description:** Shows the groups a user belongs to.
- **Common Options:**
 - `[username]`: Specify the user (default is current user).

Examples:

```
# Show current user's groups
```

```
groups
```

```
# Example output: user developers sudo
```

```
# Show another user's groups
```

```
groups newuser
```

```
# Example output: newuser : developers sudo
```

```
# Verify with id command
```

```
id newuser
```

```
# Example output: uid=1001(newuser) gid=1001(developers)
```

```
groups=1001(developers),27(sudo)
```

-
- **Note:** Groups are stored in `/etc/group`. Primary group is listed in `/etc/passwd`.

Practical Tips

- **Use `adduser` over `useradd`:** `adduser` is more user-friendly and handles home directory creation.
- **Be cautious with `sudo` group:** Adding users to `sudo` grants administrative privileges.

- **Verify changes:** Use `id` or `groups` to confirm user settings after modifications.
- **Lock accounts for security:** Use `usermod -L` or `passwd -l` to disable unused accounts.
- **Backup user files:** Copy `/etc/passwd`, `/etc/shadow`, and `/etc/group` before changes (`sudo cp /etc/passwd /etc/passwd.bak`).
- **Check `/etc/passwd` and `/etc/group`:** Use `cat` or `less` to inspect user and group configurations.

Practice Tasks

1. Create a new user `testuser` with a home directory and Bash shell.
2. Add `testuser` to the `developers` group.
3. Change the password for `testuser`.
4. Change `testuser`'s shell to `/bin/zsh`.
5. Lock the `testuser` account.
6. Verify `testuser`'s groups and account details.

Solution:

```
sudo adduser --shell /bin/bash testuser
sudo usermod -aG developers testuser
sudo passwd testuser
sudo usermod -s /bin/zsh testuser
sudo passwd -l testuser
id testuser
groups testuser
```

Linux Group Management

This document covers Linux commands for managing groups: `groupadd`, `gpasswd`, and `usermod -aG`. Each command includes a description, common options, and practical examples to help you create, modify, and manage group memberships effectively.

1. groupadd (Add a Group)

- **Description:** Creates a new group in the system, adding an entry to `/etc/group`.
- **Common Options:**
 - `-g [gid]`: Specify a group ID.
 - `-r`: Create a system group (low GID, typically for system use).
 - `-f`: Force creation, ignoring errors if the group exists.

Examples:

```
# Create a new group
```

```
sudo groupadd developers
```

```
# Create a group with a specific GID
```

```
sudo groupadd -g 2000 testers
```

```
# Create a system group
```

```
sudo groupadd -r sysgroup
```

```
# Verify group creation
```

```
cat /etc/group | grep developers
```

```
# Example output: developers:x:1001:
```

-
- **Note:** Requires `sudo`. Group details are stored in `/etc/group`. Use `getent group [groupname]` to verify.

2. gpasswd (Administer Group Password and Membership)

- **Description:** Manages group passwords and membership, including adding or removing users from a group. Rarely used for group passwords in modern systems.
- **Common Options:**
 - `-a [user]`: Add a user to the group.
 - `-d [user]`: Remove a user from the group.

- **-A [user, ...]**: Set group administrators.
- **-M [user, ...]**: Set group members (overwrites existing members).

Examples:

Add a user to a group

```
sudo gpasswd -a user1 developers
```

Remove a user from a group

```
sudo gpasswd -d user1 developers
```

Set group administrators

```
sudo gpasswd -A admin1 developers
```

Set multiple group members

```
sudo gpasswd -M user1,user2,user3 developers
```

Verify membership

```
getent group developers
```

Example output: developers:x:1001:user1,user2,user3

-
- **Note:** Requires **sudo**. Group membership changes are reflected in **/etc/group**.

3. usermod -aG (Add User to Groups)

- **Description:** Modifies a user's group membership, typically used to add a user to supplementary groups. The **-aG** combination is critical to append groups without overwriting existing ones.
- **Common Options:**
 - **-aG [group, ...]**: Append user to one or more groups.

Examples:

Add a user to multiple groups

```
sudo usermod -aG developers,testers user1
```

Verify group membership

```
groups user1
```

Example output: user1 : user1 developers testers

Check with id command

```
id user1
```

Example output: uid=1001(user1) gid=1001(user1)

groups=1001(user1),1002(developers),2000(testers)

-
- **Note:** Requires `sudo`. Always use `-a` (append) with `-G` to avoid removing existing group memberships.

Practical Tips

- **Use `groupadd` for new groups:** Specify a GID with `-g` for consistency across systems.
- **Prefer `usermod -aG` for membership:** It's simpler than `gpasswd` for adding users to groups.
- **Verify changes:** Use `groups`, `id`, or `getent group [groupname]` to confirm group memberships.
- **Backup group files:** Copy `/etc/group` before changes (`sudo cp /etc/group /etc/group.bak`).
- **System vs. regular groups:** Use `-r` with `groupadd` for system groups (e.g., for services).
- **Check `/etc/group`:** Inspect group details with `cat /etc/group` or `less /etc/group`.

Practice Tasks

1. Create a new group called `team`.
2. Add `user1` to the `team` group using `gpasswd`.
3. Add `user2` to both `team` and `developers` groups using `usermod`.
4. Remove `user1` from the `team` group using `gpasswd`.
5. Verify the group memberships of `user1` and `user2`.
6. Create a system group called `sysbackup`.

Solution:

```
sudo groupadd team
sudo gpasswd -a user1 team
sudo usermod -aG team,developers user2
sudo gpasswd -d user1 team
groups user1
groups user2
sudo groupadd -r sysbackup
```


Linux File System Basics

This document covers essential Linux commands for managing and inspecting file systems: `df`, `du`, `lsblk`, `fdisk`, and `mkfs`. Each command includes a description, common options, and practical examples to help you understand and manage file systems effectively.

1. df (Disk Free)

- **Description:** Reports disk space usage for mounted file systems.
- **Common Options:**
 - `-h`: Human-readable sizes (e.g., GB, MB).
 - `-T`: Show file system type.
 - `-i`: Show inode usage.
 - `--exclude-type=[type]`: Exclude specific file system types (e.g., `tmpfs`).

Examples:

Show disk usage in human-readable format

`df -h`

Example output: Filesystem Size Used Avail Use% Mounted on

/dev/sda1 100G 50G 50G 50% /

Show file system types

`df -T`

Example output: Includes type like ext4, xfs

Show inode usage

`df -i`

Example output: Shows inodes instead of bytes

Exclude tmpfs file systems

`df -h --exclude-type=tmpfs`

-
- **Note:** Output includes mounted file systems in `/etc/fstab` or temporary mounts.

2. du (Disk Usage)

- **Description:** Estimates disk space used by files and directories.
- **Common Options:**
 - `-h`: Human-readable sizes.
 - `-s`: Summarize total size for a directory.

- `-c`: Show grand total.
- `--max-depth=[N]`: Limit directory depth.

Examples:

Show disk usage for a directory

```
du -h /home/user
```

Example output: 4.0K /home/user/docs, 8.0M /home/user/videos

Summarize total size

```
du -sh /home/user
```

Example output: 10M /home/user

Show total with subdirectories

```
du -ch /home/user/*
```

Example output: Lists each subdirectory and total

Limit depth to 1

```
du -h --max-depth=1 /home/user
```

-
- **Note:** Use with `sudo` for restricted directories. Combine with `sort` for readability (e.g., `du -sh * | sort -h`).

3. lsblk (List Block Devices)

- **Description:** Lists block devices (e.g., disks, partitions) and their mount points.
- **Common Options:**
 - `-f`: Show file system type and mount points.
 - `-a`: Include empty devices.
 - `-o [columns]`: Specify output columns (e.g., `NAME, SIZE, FSTYPE, MOUNTPOINT`).

Examples:

List all block devices

```
lsblk
```

Example output: NAME MAJ:MIN RM SIZE TYPE MOUNTPOINT

```
# sda 8:0 0 100G disk
```

```
# └─sda1 8:1 0 100G part /
```

Show file system types and mount points

```
lsblk -f
```

Example output: Includes FSTYPE (e.g., ext4) and MOUNTPOINT

```
# Custom output columns  
lsblk -o NAME,SIZE,FSTYPE,MOUNTPOINT
```

-
- **Note:** Useful for identifying disks and partitions before formatting or mounting.

4. fdisk (Partition Table Manipulator)

- **Description:** Manages disk partitions (create, delete, modify) on block devices.
- **Common Options:**
 - `-l`: List partition tables for all or specified devices.
 - `[device]`: Specify the disk to manage (e.g., `/dev/sdb`).
- **Interactive Commands:**
 - `n`: Create new partition.
 - `d`: Delete partition.
 - `p`: Print partition table.
 - `w`: Write changes and exit.
 - `q`: Quit without saving.

Examples:

```
# List all partition tables
```

```
sudo fdisk -l
```

```
# Example output: Shows partitions for all disks
```

```
# Manage partitions on /dev/sdb
```

```
sudo fdisk /dev/sdb
```

```
# Enter interactive mode:
```

```
# n (new partition), p (primary), 1 (partition number), accept defaults
```

```
# w (write changes)
```

```
# List partitions for a specific disk
```

```
sudo fdisk -l /dev/sda
```

-
- **Note:** Requires `sudo`. Be cautious, as changes are destructive. Use `parted` for GPT disks or advanced partitioning.

5. mkfs (Make File System)

- **Description:** Creates a file system on a partition (e.g., ext4, xfs, vfat).
- **Common Commands:**
 - `mkfs.ext4 [device]`: Create an ext4 file system.

- `mkfs.vfat [device]`: Create a FAT32 file system (e.g., for USB drives).
- `-L [label]`: Set a volume label.

Examples:

Create an ext4 file system on a partition

```
sudo mkfs.ext4 /dev/sdb1
```

Warning: Erases data on /dev/sdb1

Create a FAT32 file system with a label

```
sudo mkfs.vfat -L USBDRIVE /dev/sdc1
```

Verify file system

```
lsblk -f
```

Example output: Shows new file system type

-
- **Note:** Requires `sudo`. Ensure the partition is not mounted before formatting. Use `lsblk` to identify partitions.

Practical Tips

- **Check disk usage:** Use `df -h` for a quick overview of free space.
- **Analyze large directories:** Combine `du -sh * | sort -h` to find space hogs.
- **Identify devices:** Use `lsblk -f` before partitioning or formatting.
- **Backup before partitioning:** Save data before using `fdisk` or `mkfs`, as they are destructive.
- **Use specific file systems:** Choose `ext4` for Linux, `vfat` for USB drives, or `xf`s for large storage.
- **Combine with `mount`:** After creating a file system, mount it with `mount /dev/sdb1 /mnt`.

Practice Tasks

1. Check disk usage in human-readable format.
2. Summarize the disk usage of `/home` directory.
3. List all block devices and their file system types.
4. List the partition table for `/dev/sda`.
5. Create an ext4 file system on `/dev/sdb1` (if available, use with caution).
6. Verify the new file system with `lsblk`.

Solution:

```
df -h
du -sh /home
lsblk -f
sudo fdisk -l /dev/sda
sudo mkfs.ext4 /dev/sdb1
lsblk -f
```

Linux Mounting Drives (Detailed Explanation)

This document provides an in-depth explanation of Linux commands and configurations for mounting drives: `mount`, `umount`, and `/etc/fstab`. Each section includes a detailed description, how the command works, common options, practical examples, use cases, and troubleshooting tips to help you manage file system mounts effectively.

1. mount (Mount a File System)

Overview

- **Purpose:** The `mount` command attaches a file system (e.g., on a disk partition, USB drive, or network share) to a directory in the Linux file system hierarchy, known as a *mount point*. This makes the file system's contents accessible.
- **Mechanics:** When you mount a file system, Linux maps the device's file system to a directory, allowing you to interact with it as if it were part of the main file system. The kernel handles communication between the device and the mount point.
- **Use Cases:**
 - Accessing files on a USB drive.
 - Mounting a new hard drive partition for data storage.
 - Connecting to network file systems (e.g., NFS, SMB).
- **File System Types:** Common types include `ext4` (Linux), `vfat` (FAT32 for USBs), `ntfs` (Windows), and `xf`s (large storage).

Common Options

- `-t [fstype]`: Specifies the file system type (e.g., `ext4`, `vfat`, `ntfs`). If omitted, `mount` attempts to detect it.
- `-o [options]`: Sets mount options, such as:
 - `rw`: Read-write (default for most file systems).

- **ro**: Read-only (useful for backups or damaged drives).
- **noexec**: Prevents execution of binaries.
- **nosuid**: Disables set-user-ID programs.
- **sync**: Writes changes immediately (slower but safer).
- **-L [label]**: Mounts by volume label (find with **lsblk -f**).
- **-U [uuid]**: Mounts by UUID (more reliable than device names).
- **-a**: Mounts all file systems listed in **/etc/fstab**.

Detailed Examples

1. Mounting a USB Drive:

Identify the device using **lsblk**:

lsblk

Output: NAME MAJ:MIN RM SIZE TYPE MOUNTPOINT

sdb 8:16 1 16G disk

└─sdb1 8:17 1 16G part

○

Create a mount point:

sudo mkdir /mnt/usb

○

Mount the USB drive (assuming it's **vfat**):

sudo mount -t vfat /dev/sdb1 /mnt/usb

○

Verify:

ls /mnt/usb

df -h /mnt/usb

Output: Filesystem Size Used Avail Use% Mounted on

/dev/sdb1 16G 2G 14G 13% /mnt/usb

○

2. Mounting a Partition Read-Only:

Mount an **ext4** partition read-only:

sudo mount -o ro /dev/sda2 /mnt/data

○

Verify:

```
mount | grep /mnt/data
```

```
# Output: /dev/sda2 on /mnt/data type ext4 (ro,relatime)
```

○

3. Mounting by UUID:

Find the UUID:

```
lsblk -f
```

```
# Output: NAME FSTYPE LABEL UUID
```

MOUNTPOINT

```
# sdb1 ext4 Data 1234-5678
```

○

Mount using UUID:

```
sudo mount -U 1234-5678 /mnt/data
```

○

Troubleshooting

- **Device not found:** Ensure the device exists (`lsblk` or `blkid`).
- **Wrong file system type:** Specify `-t` or check with `lsblk -f`.
- **Permission denied:** Use `sudo` or check user permissions.
- **Mount point busy:** Check for open files with `lsof /mnt/point`.

2. umount (Unmount a File System)

Overview

- **Purpose:** Detaches a file system from its mount point, making it inaccessible until remounted.
- **Mechanics:** The `umount` command tells the kernel to release the file system from the directory, ensuring all pending writes are completed. Unmounting is necessary before safely removing devices like USB drives.
- **Use Cases:**
 - Safely ejecting a USB drive.
 - Preparing a partition for reformatting.
 - Disconnecting a network share.

Common Options

- **-l**: Lazy unmount; detaches the file system when it's no longer in use (useful for busy mounts).
- **-f**: Force unmount (use with caution, may cause data loss).
- **-r**: Remount read-only if unmounting fails.

Detailed Examples

1. Unmounting a USB Drive:

Unmount the USB drive mounted at **/mnt/usb**:

```
sudo umount /mnt/usb
```

○

Verify it's unmounted:

```
mount | grep /mnt/usb
```

Should return no output

○

2. Handling a Busy Mount:

Check for processes using the mount point:

```
lsof /mnt/usb
```

Output: Lists processes accessing /mnt/usb

○

Terminate processes or use lazy unmount:

```
sudo umount -l /mnt/usb
```

○

3. Force Unmount (Emergency):

Force unmount if necessary:

```
sudo umount -f /mnt/usb
```

○

Troubleshooting

- **Device is busy**: Use **lsof** or **fuser -m /mnt/point** to find and kill processes.
- **Permission denied**: Ensure **sudo** is used.

- **Lazy unmount issues:** Check if the mount point is still referenced (`mount` command).

3. /etc/fstab (File System Table)

Overview

- **Purpose:** Defines file systems to be mounted automatically at boot or manually with `mount -a`. It centralizes mount configurations.
- **File Location:** `/etc/fstab`
- **Mechanics:** The kernel and `systemd` read `/etc/fstab` during boot to mount file systems. Each line specifies a device, mount point, and options.
- **Use Cases:**
 - Automatically mounting data drives at boot.
 - Configuring network file systems.
 - Setting up specific mount options (e.g., read-only, noexec).

File Format

```
# <device> <mount point> <fstype> <options> <dump> <pass>
UUID=1234-5678 /mnt/data ext4 defaults 0 2
```

- **Device:** Device path (e.g., `/dev/sdb1`), UUID, or label.
- **Mount Point:** Directory where the file system is mounted (e.g., `/mnt/data`).
- **Fstype:** File system type (e.g., `ext4`, `vfat`, `nfs`).
- **Options:** Mount options (e.g., `defaults`, `ro`, `noauto`, `user`).
- **Dump:** Backup flag (0 = disable, 1 = enable for `dump` utility).
- **Pass:** Filesystem check order (0 = none, 1 = root, 2 = others).

Common Options

- `defaults`: Uses default settings (rw, suid, dev, exec, auto, nouser, async).
- `noauto`: Prevents mounting at boot (requires manual `mount`).
- `user`: Allows non-root users to mount.
- `ro/rw`: Read-only or read-write.
- `noexec`: Prevents execution of binaries.

Detailed Examples

1. **Adding a New Drive to fstab:**

Find the UUID:

lsblk -f

Output: NAME FSTYPE LABEL UUID MOUNTPOINT

sdb1 ext4 Data 1234-5678

○

Create a mount point:

sudo mkdir /mnt/data

○

Edit **/etc/fstab**:

sudo nano /etc/fstab

Add:

UUID=1234-5678 /mnt/data ext4 defaults 0 2

Save and exit

○

Test the configuration:

sudo mount -a

No errors means success

○

Verify:

df -h /mnt/data

Output: Filesystem Size Used Avail Use% Mounted on

/dev/sdb1 100G 10G 90G 10% /mnt/data

○

2. Mounting a USB Drive with User Permissions:

Edit **/etc/fstab**:

sudo nano /etc/fstab

Add:

UUID=abcd-efgh /mnt/usb vfat user,noauto 0 0

Save and exit

○

Test and mount as a user:

sudo mount -a

mount /mnt/usb # As non-root user

○

Troubleshooting

- **Boot failure:** Errors in `/etc/fstab` can prevent booting. Boot into recovery mode to fix.
- **Invalid UUID:** Verify UUIDs with `blkid`.
- **Mount errors:** Test with `mount -a` and check `journalctl -b` for logs.
- **Backup fstab:** Always copy before editing (`sudo cp /etc/fstab /etc/fstab.bak`).

Practical Tips

- **Use UUIDs:** Device names (e.g., `/dev/sdb1`) can change; UUIDs are stable.
- **Test fstab changes:** Run `sudo mount -a` to catch errors before rebooting.
- **Create mount points:** Ensure mount points exist (`mkdir /mnt/something`).
- **Check processes before unmounting:** Use `lsof` or `fuser` to avoid "device busy" errors.
- **Automate mounts:** Use `/etc/fstab` for persistent mounts, `noauto` for manual control.
- **Monitor mounts:** Use `df -h` or `mount` to verify mounted file systems.

Practice Tasks

1. Identify a USB drive's UUID and mount it to `/mnt/usb` as `vfat`.
2. Unmount the USB drive and verify it's unmounted.
3. Add an entry to `/etc/fstab` to mount the USB drive at `/mnt/usb` with `user` and `noauto` options.
4. Test the `/etc/fstab` configuration.
5. Mount the USB drive as a non-root user.
6. Check the mounted drive's disk usage.

Solution:

Task 1: Identify UUID and mount

```
lsblk -f
```

Note UUID, e.g., abcd-efgh

```
sudo mkdir /mnt/usb
```

```
sudo mount -t vfat -U abcd-efgh /mnt/usb
```

Task 2: Unmount and verify

```
sudo umount /mnt/usb
```

```
mount | grep /mnt/usb
# Should return no output

# Task 3: Add to fstab
sudo nano /etc/fstab
# Add: UUID=abcd-efgh /mnt/usb vfat user,noauto 0 0
# Save and exit

# Task 4: Test fstab
sudo mount -a
# No errors means success

# Task 5: Mount as non-root user
mount /mnt/usb

# Task 6: Check disk usage
df -h /mnt/usb
```

Linux Basic LVM Concepts

This document provides an in-depth explanation of Logical Volume Management (LVM) in Linux, focusing on the commands `pvcreate`, `vgcreate`, and `lvcreate`. It includes detailed descriptions, how LVM works, common options, practical examples, use cases, and troubleshooting tips to help you manage flexible storage effectively.

Overview of LVM

- **Purpose:** Logical Volume Management (LVM) provides a flexible way to manage disk storage by abstracting physical disks into logical volumes. It allows resizing, snapshotting, and combining multiple disks into a single logical storage pool.
- **Components:**
 - **Physical Volumes (PVs):** Physical disks or partitions initialized for LVM (created with `pvcreate`).
 - **Volume Groups (VGs):** A pool of storage combining one or more PVs (created with `vgcreate`).
 - **Logical Volumes (LVs):** Virtual partitions carved out of a VG, which can be formatted and mounted (created with `lvcreate`).
- **Use Cases:**
 - Resizing file systems without reformatting.
 - Combining multiple disks into a single logical volume.

- Creating snapshots for backups.
- **Advantages:**
 - Dynamic resizing of volumes.
 - Spanning storage across multiple disks.
 - Simplified storage management.
- **Mechanics:** LVM uses a layered approach where PVs are grouped into VGs, and LVs are allocated from VGs. LVs can be formatted with file systems (e.g., `ext4`) and mounted.

1. pvcreate (Create Physical Volume)

Overview

- **Purpose:** Initializes a disk or partition as a Physical Volume (PV) for use in LVM.
- **Mechanics:** Marks the device with LVM metadata, making it available to add to a Volume Group.
- **Use Cases:**
 - Preparing a new disk for LVM.
 - Adding an existing partition to an LVM setup.

Common Options

- `--force`: Force creation, overwriting existing data (use with caution).
- `--setphysicalvolumesize [size]`: Set the usable size of the PV.

Detailed Examples

1. Initialize a Disk as a PV:

Identify the disk (ensure it's unmounted and has no critical data):

```
lsblk
```

```
# Output: NAME MAJ:MIN RM SIZE TYPE MOUNTPOINT
```

```
# sdb 8:16 0 200G disk
```

○

Create the PV:

```
sudo pvcreate /dev/sdb
```

```
# Output: Physical volume "/dev/sdb" successfully created.
```

○

Verify:

`pvdisplay`

Output: Shows PV details like size and status

○

2. Initialize a Partition:

Create a partition (e.g., using `fdisk`):

`sudo fdisk /dev/sdb`

Create new partition (n), set type to LVM (t, 8e), write (w)

○

Initialize the partition:

`sudo pvcreate /dev/sdb1`

○

Troubleshooting

- **Device in use:** Ensure the disk/partition is not mounted (`umount /dev/sdb1`) or used by another process (`lsof /dev/sdb`).
- **Invalid device:** Verify the device exists (`lsblk`).
- **Data loss warning:** `pvcreate` wipes data; back up before use.

2. vgcreate (Create Volume Group)

Overview

- **Purpose:** Combines one or more Physical Volumes into a Volume Group (VG), creating a storage pool for Logical Volumes.
- **Mechanics:** Aggregates PVs into a single manageable unit, allowing allocation of Logical Volumes.
- **Use Cases:**
 - Grouping multiple disks for unified storage.
 - Preparing for Logical Volume creation.

Common Options

- `-s [size]`: Set physical extent size (default is 4MiB).
- `--name [vgname]`: Specify the Volume Group name.

Detailed Examples

1. Create a Volume Group:

Assume `/dev/sdb` is a PV:

```
sudo vgcreate myvg /dev/sdb
```

Output: Volume group "myvg" successfully created

○

Verify:

```
vgdisplay
```

Output: Shows VG name, size, and available space

○

2. Add Multiple PVs:

Initialize additional PVs:

```
sudo pvcreate /dev/sdc
```

○

Create VG with multiple PVs:

```
sudo vgcreate myvg /dev/sdb /dev/sdc
```

○

Troubleshooting

- **PV not found:** Ensure PVs are initialized (`pvdisplay`).
- **VG name exists:** Use a unique name or remove the existing VG (`vgremove`).
- **Insufficient space:** Check PV sizes (`pvdisplay`).

3. lvcreate (Create Logical Volume)

Overview

- **Purpose:** Allocates a Logical Volume (LV) from a Volume Group, which can be formatted and mounted like a regular partition.
- **Mechanics:** Carves out a portion of the VG's storage for the LV, which can then be formatted with a file system (e.g., `ext4`).
- **Use Cases:**

- Creating resizable storage for data.
- Setting up volumes for specific applications.

Common Options

- **-L [size]**: Specify the size of the LV (e.g., **10G** for 10GB).
- **-n [name]**: Name the Logical Volume.
- **-l [extents]**: Specify size in extents (e.g., **50%FREE** for 50% of available VG space).

Detailed Examples

1. Create a Logical Volume:

Create an LV of 10GB in **myvg**:

```
sudo lvcreate -L 10G -n mylv myvg
```

Output: Logical volume "mylv" created

○

Verify:

```
lvdisplay
```

Output: Shows LV path (e.g., /dev/myvg/mylv)

○

2. Format and Mount the LV:

Format the LV as **ext4**:

```
sudo mkfs.ext4 /dev/myvg/mylv
```

○

Create a mount point and mount:

```
sudo mkdir /mnt/mylv
```

```
sudo mount /dev/myvg/mylv /mnt/mylv
```

○

Verify:

```
df -h /mnt/mylv
```

Output: Filesystem Size Used Avail Use% Mounted on

```
# /dev/myvg/mylv      10G   0G   10G   0%   /mnt/mylv
```

○

3. Create an LV Using All Free Space:

Use all available VG space:

```
sudo lvcreate -l 100%FREE -n data1v myvg
```

○

Troubleshooting

- **Insufficient VG space:** Check available space (`vgdisplay`).
- **LV path issues:** Use `lvdisplay` to confirm the LV path (e.g., `/dev/vgname/lvname`).
- **Mount failures:** Ensure the LV is formatted (`mkfs`) and the mount point exists.

Practical Example: Setting Up an LVM Volume

Scenario

You have a new 200GB disk (`/dev/sdb`) and want to create a 50GB `ext4` logical volume for data storage.

Steps

Initialize the Disk as a PV:

```
sudo pvcreate /dev/sdb  
pvdisplay
```

1.

Create a Volume Group:

```
sudo vgcreate datavg /dev/sdb  
vgdisplay
```

2.

Create a Logical Volume:

```
sudo lvcreate -L 50G -n data1v datavg  
lvdisplay
```

3.

Format and Mount the LV:

```
sudo mkfs.ext4 /dev/datavg/data1v
sudo mkdir /mnt/data
sudo mount /dev/datavg/data1v /mnt/data
df -h /mnt/data
```

4.

5. **Add to /etc/fstab for Persistence:**

Find the LV's UUID:

```
lsblk -f
```

```
# Output: NAME          FSTYPE LABEL UUID          MOUNTPOINT
# datavg-data1v  ext4    abcd-1234  /mnt/data
```

○

Edit **/etc/fstab**:

```
sudo nano /etc/fstab
# Add: UUID=abcd-1234 /mnt/data ext4 defaults 0 2
# Save and exit
```

○

Test:

```
sudo mount -a
```

○

Practical Tips

- **Backup data:** LVM operations like **pvcreeate** and **mkfs** are destructive; back up critical data.
- **Use UUIDs in fstab:** LVs have stable UUIDs (**lsblk -f** or **blkid**).
- **Check LVM status:** Use **pvs**, **vgs**, and **lvs** for quick summaries of PVs, VGs, and LVs.
- **Extend LVs:** Use **lvextend** to resize LVs (e.g., **sudo lvextend -L +10G /dev/myvg/mylv**).
- **Monitor space:** Use **vgdisplay** to check free space in a VG before creating LVs.
- **Install LVM:** Ensure **lvm2** is installed (**sudo apt install lvm2**).

Practice Tasks

1. Initialize a disk (`/dev/sdc`) as a Physical Volume.
2. Create a Volume Group named `testvg` using `/dev/sdc`.
3. Create a 20GB Logical Volume named `testlv` in `testvg`.
4. Format `testlv` as `ext4` and mount it to `/mnt/test`.
5. Add `testlv` to `/etc/fstab` for automatic mounting.
6. Verify the mount and check disk usage.

Solution:

```
sudo pvcreate /dev/sdc
sudo vgcreate testvg /dev/sdc
sudo lvcreate -L 20G -n testlv testvg
sudo mkfs.ext4 /dev/testvg/testlv
sudo mkdir /mnt/test
sudo mount /dev/testvg/testlv /mnt/test
lsblk -f
# Note UUID, e.g., 5678-abcd
sudo nano /etc/fstab
# Add: UUID=5678-abcd /mnt/test ext4 defaults 0 2
# Save and exit
sudo mount -a
df -h /mnt/test
```

Linux File Permissions

This document provides an in-depth explanation of Linux commands for managing file permissions: `chmod`, `chown`, `setfacl`, and `getfacl`. Each section includes a detailed description, how the command works, common options, practical examples, use cases, and troubleshooting tips to help you manage file and directory access effectively.

Overview of File Permissions

- **Purpose:** Linux file permissions control who can read, write, or execute files and directories. Permissions are managed for three categories: owner, group, and others.
- **Mechanics:**
 - **Standard Permissions:** Defined by read (`r`), write (`w`), and execute (`x`) for owner, group, and others, represented as `rw-rw-rw-` (9 characters) or octal (e.g., `755`).
 - **Access Control Lists (ACLs):** Extend standard permissions with fine-grained control using `setfacl` and `getfacl`.

- **Ownership:** Each file has an owner (user) and a group, managed by `chown`.
- **Use Cases:**
 - Securing sensitive files (e.g., configuration files).
 - Granting specific users access to shared directories.
 - Controlling script execution permissions.
- **Key Files:**
 - Permissions and ownership are stored in the file system's inode metadata.
 - View with `ls -l` (e.g., `-rw-r--r-- user group`).

1. chmod (Change Mode)

Overview

- **Purpose:** Modifies file or directory permissions for owner, group, and others.
- **Mechanics:** Uses symbolic (e.g., `u+x`) or octal (e.g., `755`) notation to set permissions.
 - **Symbolic:** `u` (user/owner), `g` (group), `o` (others), `a` (all); `+` (add), `-` (remove), `=` (set); `r` (read), `w` (write), `x` (execute).
 - **Octal:** Three digits (0-7) for owner, group, others (e.g., `755 = rwxr-xr-x`).
 - `4` = read, `2` = write, `1` = execute (sum them: `7 = rwx`, `6 = rw-`, etc.).
- **Use Cases:**
 - Making scripts executable.
 - Restricting access to sensitive files.
 - Setting directory permissions for collaboration.

Common Options

- `-R`: Apply permissions recursively to directories and their contents.
- `-v`: Verbose output, showing changed files.
- `--reference=[file]`: Copy permissions from another file.

Detailed Examples

1. Set Permissions with Octal Notation:

Make a script executable by owner, read-only for group and others:

```
chmod 744 script.sh
```

```
ls -l script.sh
```

```
# Output: -rwxr--r-- 1 user group 123 Jun 27 2025 script.sh
```

○

2. Set Permissions with Symbolic Notation:

Add execute permission for owner:

```
chmod u+x script.sh
```

```
ls -l script.sh
```

```
# Output: -rwxr--r-- 1 user group 123 Jun 27 2025 script.sh
```

○

3. Recursive Permissions for a Directory:

Set a directory to **rwxr-x---** (owner full access, group read/execute, others none):

```
chmod -R 750 /home/user/docs
```

```
ls -ld /home/user/docs
```

```
# Output: drwxr-x--- 2 user group 4096 Jun 27 2025 docs
```

○

Troubleshooting

- **Permission denied:** Ensure you have **sudo** or ownership to change permissions.
- **Recursive errors:** Use **chmod -R** carefully; verify with **ls -lR**.
- **Incorrect mode:** Check octal (**0-7**) or symbolic syntax (**u, g, o**).

2. chown (Change Owner)

Overview

- **Purpose:** Changes the owner and/or group of a file or directory.
- **Mechanics:** Updates the file's inode metadata to reflect new user and/or group ownership, referencing **/etc/passwd** and **/etc/group**.
- **Use Cases:**
 - Transferring file ownership to another user.
 - Assigning files to a group for shared access.
 - Fixing ownership after copying files.

Common Options

- **-R:** Apply changes recursively.
- **-v:** Verbose output.
- **--reference=[file]:** Copy ownership from another file.

Detailed Examples

1. Change File Owner:

Change owner of a file to **newuser**:

```
sudo chown newuser file.txt
```

```
ls -l file.txt
```

```
# Output: -rw-r--r-- 1 newuser group 123 Jun 27 2025 file.txt
```

○

2. Change Owner and Group:

Change owner to **newuser** and group to **developers**:

```
sudo chown newuser:developers file.txt
```

```
ls -l file.txt
```

```
# Output: -rw-r--r-- 1 newuser developers 123 Jun 27 2025 file.txt
```

○

3. Recursive Ownership Change:

Change ownership of a directory and its contents:

```
sudo chown -R user1:users /home/user1/docs
```

```
ls -lR /home/user1/docs
```

○

Troubleshooting

- **User/group not found:** Verify users (**id user**) and groups (**getent group groupname**).
- **Permission denied:** Use **sudo** or ensure you have sufficient privileges.
- **Recursive issues:** Confirm the directory path and use **ls -lR** to verify changes.

3. setfacl (Set File Access Control Lists)

Overview

- **Purpose:** Sets Access Control Lists (ACLs) to provide fine-grained permissions beyond standard owner/group/others.
- **Mechanics:** ACLs extend permissions by defining specific access for users or groups, stored in the file system's extended attributes.

- **Use Cases:**
 - Granting specific users access to a file without changing group ownership.
 - Allowing multiple groups different permissions on a directory.
 - Managing shared directories with complex access rules.

Common Options

- **-m [acl]**: Modify or add ACL entries (e.g., **u:user:rw**, **g:group:rx**).
- **-x [acl]**: Remove an ACL entry.
- **-R**: Apply recursively.
- **-b**: Remove all ACLs, reverting to standard permissions.
- **--set**: Set new ACLs, replacing existing ones.

Detailed Examples

1. Grant a User Specific Permissions:

Allow **user2** read and write access to a file:

```
setfacl -m u:user2:rw file.txt
getfacl file.txt
# Output: user::rw-
#      group::r--
#      other::r--
#      user:user2:rw-
```

○

2. Grant a Group Permissions:

Allow **developers** group read and execute access:

```
setfacl -m g:developers:rx /home/user/docs
getfacl /home/user/docs
# Output: Includes group:developers:r-x
```

○

3. Recursive ACLs:

Apply ACLs to a directory and its contents:

```
setfacl -R -m u:user2:rwx /home/user/docs
```

○

Troubleshooting

- **ACL not supported:** Ensure the file system supports ACLs (e.g., `ext4`, `xfs`; enable with `mount -o acl`).
- **Command not found:** Install `acl` package (`sudo apt install acl`).
- **Unexpected access:** Verify ACLs with `getfacl` and check standard permissions (`ls -l`).

4. getfacl (Get File Access Control Lists)

Overview

- **Purpose:** Displays the ACLs of a file or directory, including standard permissions and extended ACL entries.
- **Mechanics:** Reads extended attributes to show detailed access rules.
- **Use Cases:**
 - Verifying ACLs after modification.
 - Debugging permission issues.
 - Documenting access rules for shared resources.

Common Options

- `-R`: Display ACLs recursively for directories.
- `--omit-header`: Skip header information for cleaner output.

Detailed Examples

View ACLs for a File:

```
getfacl file.txt
# Output: # file: file.txt
# owner: user
# group: group
user::rw-
user:user2:rw-
group::r--
other::r--
```

1.

View ACLs Recursively:

```
getfacl -R /home/user/docs
```


Output: Shows ACLs for all files and subdirectories

2.

Clean Output:

```
getfacl --omit-header file.txt
# Output: user::rw-
#   user:user2:rw-
#   group::r--
#   other::r--
```

3.

Troubleshooting

- **No ACLs displayed:** File system may not support ACLs or none are set.
- **Permission denied:** Use `sudo` for restricted files.
- **Complex output:** Use `ls -l` alongside `getfacl` to compare standard permissions.

Practical Tips

- **Check permissions first:** Use `ls -l` to view standard permissions before modifying.
- **Use octal for `chmod`:** Octal notation (e.g., `755`) is faster for setting all permissions at once.
- **Secure sensitive files:** Use `600` for private files and `700` for directories.
- **Combine `chown` and `chmod`:** Set ownership and permissions together for new files.
- **Use ACLs sparingly:** Reserve ACLs for complex permission scenarios; standard permissions are simpler.
- **Backup before changes:** Copy critical files or directories before modifying permissions.

Practice Tasks

1. Create a file `test.txt` and set permissions to `rw-r-----` using `chmod`.
2. Change the owner of `test.txt` to `user1` and group to `developers`.
3. Grant `user2` read and execute permissions on `test.txt` using `setfacl`.
4. Verify the ACLs on `test.txt` with `getfacl`.
5. Recursively set permissions on `/home/user/docs` to `rwxr-x---` and add `user2` with full access via ACLs.
6. Check ownership and permissions of `/home/user/docs`.

Solution:

```
touch test.txt
chmod 640 test.txt
sudo chown user1:developers test.txt
setfacl -m u:user2:rx test.txt
getfacl test.txt
chmod -R 750 /home/user/docs
setfacl -R -m u:user2:rwX /home/user/docs
ls -ld /home/user/docs
getfacl /home/user/docs
```

Linux Logs and Log Management

This document provides an in-depth explanation of Linux tools and configurations for managing system logs: `tail`, `grep`, `/var/log`, and `logrotate`. Each section includes a detailed description, how the tool or configuration works, common options, practical examples, use cases, and troubleshooting tips to help you monitor and manage logs effectively.

Overview of Logs

- **Purpose:** Logs record system and application events, such as errors, warnings, authentication attempts, and service activities, aiding in debugging, monitoring, and auditing.
- **Mechanics:** Logs are typically stored as text files in `/var/log`, managed by the system logger (e.g., `rsyslog` or `syslog`). Tools like `tail` and `grep` help analyze logs, while `logrotate` manages log file growth.
- **Use Cases:**
 - Troubleshooting service failures (e.g., checking `nginx` logs).
 - Monitoring security events (e.g., failed login attempts).
 - Managing disk space by rotating old logs.
- **Key Directory:** `/var/log` contains most system and service logs (e.g., `/var/log/syslog`, `/var/log/auth.log`).

1. tail (View End of File)

Overview

- **Purpose:** Displays the last few lines of a file, useful for monitoring log files in real-time or checking recent events.

- **Mechanics:** Reads the end of a file (default 10 lines) and can follow changes as they occur.
- **Use Cases:**
 - Monitoring live logs (e.g., web server errors).
 - Checking the most recent entries in a log file.

Common Options

- `-n [lines]`: Show the last `n` lines.
- `-f`: Follow the file, displaying new lines as they are written.
- `-q`: Quiet mode (suppress headers for multiple files).
- `--pid=[pid]`: Stop following when a specific process ends.

Detailed Examples

View Last 10 Lines of a Log:

```
tail /var/log/syslog
# Output: Shows the last 10 lines of system logs
```

1.

View Last 20 Lines:

```
tail -n 20 /var/log/auth.log
# Output: Shows the last 20 lines of authentication logs
```

2.

Monitor Logs in Real-Time:

```
tail -f /var/log/nginx/error.log
# Output: Displays new error log entries as they occur
# Press Ctrl+C to stop
```

3.

Follow Until Process Ends:

```
# Find nginx PID
pidof nginx
# Output: e.g., 1234
tail -f --pid=1234 /var/log/nginx/access.log
# Stops when nginx process ends
```

4.

Troubleshooting

- **Permission denied:** Use `sudo` for restricted logs (e.g., `sudo tail /var/log/auth.log`).
- **No output:** Verify the log file exists (`ls /var/log`) and has recent entries.
- **Too much output:** Use `-n` to limit lines or combine with `grep`.

2. grep (Search Text in Files)

Overview

- **Purpose:** Searches for patterns in log files or other text, filtering relevant lines for analysis.
- **Mechanics:** Uses regular expressions to match text, outputting matching lines or counts.
- **Use Cases:**
 - Finding error messages in logs.
 - Filtering logs by user or IP address.
 - Counting occurrences of specific events.

Common Options

- `-i`: Case-insensitive search.
- `-r`: Recursive search in directories.
- `-n`: Show line numbers.
- `-c`: Count matching lines.
- `-v`: Show non-matching lines (inverse search).
- `-w`: Match whole words only.

Detailed Examples

Search for Errors in a Log:

```
grep "error" /var/log/syslog  
# Output: Lines containing "error" (case-sensitive)
```

1.

Case-Insensitive Search with Line Numbers:

```
grep -in "fail" /var/log/auth.log  
# Output: e.g., 123:Failed password for user1
```

2.

Count Failed Logins:

```
grep -c "Failed password" /var/log/auth.log
```

Output: e.g., 5

3.

Recursive Search in Log Directory:

```
grep -r "crash" /var/log
```

Output: Matches in all files under /var/log

4.

Troubleshooting

- **No matches:** Check case sensitivity (`-i`) or pattern accuracy.
- **Permission denied:** Use `sudo` for restricted files.
- **Too many matches:** Narrow the search with `-w` or combine with `tail`.

3. /var/log (System Log Directory)

Overview

- **Purpose:** Stores log files for the system and services, managed by the system logger (e.g., `rsyslog` or `systemd-journald`).
- **Mechanics:** Log files are written by services, daemons, and the kernel. Common files include:
 - `/var/log/syslog` or `/var/log/messages`: General system logs.
 - `/var/log/auth.log`: Authentication-related logs.
 - `/var/log/kern.log`: Kernel logs.
 - Service-specific logs (e.g., `/var/log/nginx/access.log`).
- **Use Cases:**
 - Debugging system or application issues.
 - Auditing security events.
 - Monitoring service activity.

Common Log Files

- `/var/log/syslog` (Ubuntu) or `/var/log/messages` (CentOS): System-wide logs.
- `/var/log/auth.log` (Ubuntu) or `/var/log/secure` (CentOS): Authentication logs.

- `/var/log/dmesg`: Kernel ring buffer (boot and hardware messages).
- `/var/log/nginx/*.log`: Web server logs (if nginx is installed).

Detailed Examples

View System Logs:

```
sudo less /var/log/syslog  
# Navigate with arrow keys, q to quit
```

1.

Check Recent Authentication Events:

```
sudo tail -n 10 /var/log/auth.log  
# Output: Last 10 authentication events
```

2.

Search for Kernel Issues:

```
sudo grep "error" /var/log/kern.log
```

3.

Combine with journalctl:

```
journalctl --since "2025-06-27" -u nginx  
# Output: nginx logs from systemd journal
```

4.

Troubleshooting

- **File not found**: Ensure the log file exists (`ls /var/log`) and the service is running.
- **Empty logs**: Check if the service is configured to log or if `rsyslog/systemd-journald` is active.
- **Permission issues**: Use `sudo` to access restricted logs.

4. logrotate (Rotate Log Files)

Overview

- **Purpose**: Manages log file growth by rotating, compressing, and deleting old logs to save disk space.

- **Mechanics:** Controlled by configuration files in `/etc/logrotate.conf` and `/etc/logrotate.d/`. Runs periodically via cron (`/etc/cron.daily/logrotate`).
- **Use Cases:**
 - Preventing logs from filling disk space.
 - Archiving old logs for auditing.
 - Compressing logs to save space.

Configuration File

- **Location:** `/etc/logrotate.conf` (global settings), `/etc/logrotate.d/` (service-specific configs).

Format:

```
/var/log/example.log {
    daily
    rotate 7
    compress
    missingok
    notifempty
}
```

- - **Options:**
 - `daily/weekly/monthly`: Rotation frequency.
 - `rotate [n]`: Keep `n` old logs.
 - `compress`: Compress old logs (e.g., `.gz`).
 - `missingok`: Ignore missing log files.
 - `notifempty`: Skip rotation if the log is empty.

Detailed Examples

View Global Logrotate Config:

```
cat /etc/logrotate.conf
# Output: Shows default rotation settings
```

1.

Check Service-Specific Config:

```
cat /etc/logrotate.d/nginx
# Example output:
# /var/log/nginx/*.log {
#     daily
```

```
# rotate 14
# compress
# missingok
# }
```

2.

Manually Run Logrotate:

```
sudo logrotate -f /etc/logrotate.conf
# Forces rotation of all logs
```

3.

Create a Custom Logrotate Config:

```
sudo nano /etc/logrotate.d/myapp
# Add:
/var/log/myapp.log {
    weekly
    rotate 4
    compress
    missingok
}
# Save and exit
# Test:
sudo logrotate -d /etc/logrotate.d/myapp
# Debug mode, shows what would happen
```

4.

Troubleshooting

- **Logs not rotating:** Check cron job (`/etc/cron.daily/logrotate`) and permissions.
- **Config errors:** Test with `logrotate -d` to debug.
- **Disk space issues:** Verify old logs are compressed or deleted (`ls /var/log`).

Practical Tips

- **Monitor logs in real-time:** Use `tail -f` for live debugging.
- **Filter efficiently:** Combine `grep` with `tail` (e.g., `tail -f /var/log/syslog | grep error`).

- **Secure log files:** Ensure logs like `/var/log/auth.log` have restricted permissions (`ls -l /var/log`).
- **Check disk usage:** Use `du -sh /var/log/*` to monitor log sizes.
- **Automate log analysis:** Script `grep` searches in cron jobs for alerts.
- **Backup configs:** Copy `/etc/logrotate.conf` and `/etc/logrotate.d/` before editing.

Practice Tasks

1. View the last 15 lines of `/var/log/syslog`.
2. Search for "error" in `/var/log/nginx/error.log`.
3. Monitor `/var/log/auth.log` in real-time for failed logins.
4. Create a `logrotate` config for `/var/log/myapp.log` to rotate weekly, keeping 5 logs.
5. Test the `logrotate` config in debug mode.
6. Check the size of all logs in `/var/log`.

Solution:

```

sudo tail -n 15 /var/log/syslog
sudo grep "error" /var/log/nginx/error.log
sudo tail -f /var/log/auth.log | grep "Failed password"
sudo nano /etc/logrotate.d/myapp
# Add:
# /var/log/myapp.log {
#   weekly
#   rotate 5
#   compress
#   missingok
# }
# Save and exit
sudo logrotate -d /etc/logrotate.d/myapp
du -sh /var/log/*

```

Linux Process Management

This document provides an in-depth explanation of Linux commands for managing processes: `ps`, `top`, `htop`, `kill`, and `nice`. Each section includes a detailed description, how the command works, common options, practical examples, use cases, and troubleshooting tips to help you monitor and control processes effectively.

Overview of Process Management

- **Purpose:** Process management involves monitoring, controlling, and prioritizing running processes (programs in execution) on a Linux system.
- **Mechanics:** Each process has a unique Process ID (PID), user, priority, and resource usage. The kernel manages processes, and tools like `ps`, `top`, and `htop` provide visibility, while `kill` and `nice` control execution.
- **Use Cases:**
 - Identifying resource-heavy processes.
 - Terminating unresponsive applications.
 - Adjusting process priorities for performance.
- **Key Concepts:**
 - **PID:** Unique identifier for a process.
 - **PPID:** Parent Process ID (the process that started it).
 - **Priority:** Determines CPU scheduling (lower `nice` value = higher priority).

1. ps (Process Status)

Overview

- **Purpose:** Displays a snapshot of current processes, including their PID, user, and command.
- **Mechanics:** Queries the `/proc` file system to show process details.
- **Use Cases:**
 - Listing all processes for a user.
 - Finding the PID of a specific process.
 - Debugging process ownership or status.

Common Options

- `-e` or `aux`: Show all processes (BSD style: `aux`).
- `-u [user]`: Show processes for a specific user.
- `-f`: Full format, including PPID and command arguments.
- `--forest`: Show process hierarchy (tree view).

- `-p [pid]`: Show specific PID(s).

Detailed Examples

List All Processes:

```
ps aux
# Output: USER PID %CPU %MEM  VSZ  RSS TTY STAT START TIME COMMAND
#      user1 1234  0.1  0.5 123456 7890 ?   S   10:00 0:01 /usr/bin/nginx
```

1.

Show Processes for a User:

```
ps -u user1
# Output: Lists processes owned by user1
```

2.

Display Process Hierarchy:

```
ps --forest
# Output: Shows parent-child process relationships
```

3.

Find a Specific Process:

```
ps aux | grep nginx
# Output: Lists processes matching "nginx"
```

4.

Troubleshooting

- **No output:** Ensure the process exists or use `ps -e` for all processes.
- **Too much output:** Pipe to `grep` (e.g., `ps aux | grep [pattern]`).
- **Missing details:** Use `-f` for full output or specify columns with `-o`.

2. top (Table of Processes)

Overview

- **Purpose:** Provides a real-time, interactive view of system processes, including CPU, memory usage, and process details.

- **Mechanics:** Continuously updates process information, allowing sorting and filtering.
- **Use Cases:**
 - Monitoring system resource usage.
 - Identifying high-CPU or high-memory processes.
 - Interactively killing processes.

Common Keys (Interactive)

- **q:** Quit.
- **k:** Kill a process (prompts for PID).
- **r:** Renice a process (change priority).
- **f:** Manage fields (add/remove columns).
- **P:** Sort by CPU usage, **M:** Sort by memory.

Detailed Examples

Start top:

top
Output: Real-time display with CPU, memory, and process list

1.

Kill a Process in top:

top
Press k, enter PID (e.g., 1234), confirm with Enter

2.

Sort by Memory Usage:

top
Press M to sort by %MEM

3.

Troubleshooting

- **High CPU load:** Check **%CPU** column to identify culprits.
- **Unresponsive:** Use **q** to exit or **Ctrl+C** if frozen.
- **Missing processes:** Ensure you have permissions (**sudo top** for all processes).

3. htop (Enhanced top)

Overview

- **Purpose:** A user-friendly, interactive process viewer with a colorful interface, mouse support, and scrolling.
- **Mechanics:** Similar to `top` but with improved visuals and usability, built on `ncurses`.
- **Use Cases:**
 - Easier process monitoring with intuitive navigation.
 - Filtering or searching processes by name.
 - Adjusting priorities or killing processes interactively.

Common Keys (Interactive)

- `q` or `F10`: Quit.
- `F9`: Kill a process.
- `F7/F8`: Decrease/increase nice value (priority).
- `F3`: Search for a process.
- `F5`: Show process tree.

Detailed Examples

Start htop:

```
htop
# Output: Colorful interface with process list, CPU/memory bars
```

1.

Search for a Process:

```
htop
# Press F3, type "nginx", Enter
# Highlights matching processes
```

2.

Kill a Process:

```
htop
# Select process with arrow keys, press F9, choose signal (e.g., SIGTERM), Enter
```

3.

Troubleshooting

- **Command not found:** Install `htop` (`sudo apt install htop`).

- **Permission issues:** Use `sudo htop` for full process visibility.
- **Navigation issues:** Use arrow keys or mouse for selection.

4. kill (Send Signal to Process)

Overview

- **Purpose:** Sends signals to processes to terminate or control them (e.g., SIGTERM for graceful stop, SIGKILL for forced termination).
- **Mechanics:** Uses the kernel's signal system to communicate with processes, identified by PID.
- **Use Cases:**
 - Stopping unresponsive applications.
 - Sending custom signals to daemons (e.g., reload configuration).
- **Common Signals:**
 - `1` (SIGHUP): Reload configuration.
 - `9` (SIGKILL): Force kill (non-graceful).
 - `15` (SIGTERM): Graceful termination (default).

Common Options

- `-s [signal]`: Specify signal by name or number.
- `-l`: List available signals.

Detailed Examples

Gracefully Terminate a Process:

```
# Find PID
ps aux | grep nginx
# Output: user1 1234 ... /usr/bin/nginx
sudo kill 1234
# Sends SIGTERM
```

1.

Force Kill a Process:

```
sudo kill -9 1234
# Sends SIGKILL
```

2.

List Signals:

kill -l

Output: Lists signals like 1) SIGHUP, 9) SIGKILL, 15) SIGTERM

3.

Troubleshooting

- **Process not terminated:** Use **SIGKILL** (-9) if **SIGTERM** fails.
- **Permission denied:** Use **sudo** for processes owned by other users.
- **Wrong PID:** Verify PID with **ps** or **pidof**.

5. nice (Set Process Priority)

Overview

- **Purpose:** Sets the priority (niceness) of a new process or adjusts existing ones with **renice**. Lower nice values = higher priority.
- **Mechanics:** Niceness ranges from -20 (highest priority) to 19 (lowest). Default is 0.
- **Use Cases:**
 - Prioritizing critical tasks (e.g., database backups).
 - Reducing priority for background tasks (e.g., file indexing).

Common Options

- **-n [value]:** Set nice value (e.g., **-n 10** for lower priority).
- For **renice**: **[value] -p [pid]** to adjust an existing process.

Detailed Examples

Run a Command with Low Priority:

```
nice -n 10 tar -czf backup.tar.gz /home/user
```

Runs tar with lower priority

1.

Adjust Priority of a Running Process:

```
# Find PID
```

```
ps aux | grep nginx
```

```
# Output: user1 1234 ... /usr/bin/nginx
```

```
sudo renice 5 -p 1234
```

Sets nice value to 5

2.

Run with High Priority:

```
sudo nice -n -10 /usr/bin/mysql
```

Requires sudo for negative values

3.

Troubleshooting

- **Permission denied:** Use `sudo` for negative nice values or other users' processes.
- **No effect:** Check current priority (`ps -l` or `top`) and system load.
- **Invalid value:** Nice values must be between -20 and 19.

Practical Tips

- **Use `htop` for ease:** It's more intuitive than `top` for monitoring and managing processes.
- **Filter processes:** Combine `ps` with `grep` (e.g., `ps aux | grep [pattern]`).
- **Be cautious with SIGKILL:** Use `kill -9` only as a last resort, as it prevents graceful shutdown.
- **Monitor resources:** Use `top` or `htop` to identify resource-heavy processes before adjusting priorities.
- **Check process status:** Use `ps` or `top` to verify a process is running after starting.
- **Automate monitoring:** Script process checks with `ps` and `grep` in cron jobs.

Practice Tasks

1. List all processes owned by `user1` using `ps`.
2. Monitor system processes in real-time with `top`.
3. Use `htop` to find and kill a process named `sleep`.
4. Find the PID of `nginx` and terminate it gracefully.
5. Run a `tar` command with low priority (nice value 15).
6. Adjust the priority of an existing process (PID 1234) to 10.

Solution:

```
ps -u user1
```

```
top
```

```
htop
```



```
# In htop: F3, type "sleep", F9, select SIGTERM, Enter
ps aux | grep nginx
# Note PID, e.g., 1234
sudo kill 1234
nice -n 15 tar -czf backup.tar.gz /home/user
sudo renice 10 -p 1234
```

Linux Package Management

This document provides an in-depth explanation of Linux commands for managing packages: `apt`, `yum/dnf`, `dpkg`, and `rpm`. Each section includes a detailed description, how the tool works, common options, practical examples, use cases, and troubleshooting tips to help you install, update, and manage software packages effectively.

Overview of Package Management

- **Purpose:** Package management tools handle the installation, updating, and removal of software packages, including dependencies, on Linux systems.
- **Mechanics:** Packages are precompiled software bundles (e.g., `.deb` for Debian-based systems, `.rpm` for Red Hat-based systems) stored in repositories. High-level tools like `apt` and `yum/dnf` manage repositories, while low-level tools like `dpkg` and `rpm` handle individual package files.
- **Use Cases:**
 - Installing software (e.g., web servers, editors).
 - Updating system packages for security and performance.
 - Removing unused software to free disk space.
- **Key Distributions:**
 - **Debian-based** (e.g., Ubuntu): Use `apt` and `dpkg`.
 - **Red Hat-based** (e.g., CentOS, Fedora): Use `yum` (older) or `dnf` (newer) and `rpm`.

1. apt (Advanced Package Tool)

Overview

- **Purpose:** A high-level package manager for Debian-based systems (e.g., Ubuntu) that handles package installation, updates, and dependency resolution.
- **Mechanics:** Communicates with repositories defined in `/etc/apt/sources.list` or `/etc/apt/sources.list.d/` to download and install `.deb` packages.

- **Use Cases:**
 - Installing tools like `nginx` or `vim`.
 - Updating all system packages.
 - Managing repository configurations.

Common Options

- `update`: Refresh repository metadata.
- `upgrade`: Update installed packages.
- `install [package]`: Install a package and its dependencies.
- `remove [package]`: Remove a package (leaves config files).
- `purge [package]`: Remove a package and its config files.
- `autoremove`: Remove unneeded dependencies.
- `-y`: Assume "yes" to prompts.

Detailed Examples

Update Package Lists:

```
sudo apt update
```

Output: Fetches latest package metadata from repositories

1.

Upgrade Installed Packages:

```
sudo apt upgrade
```

Output: Updates all installed packages to latest versions

2.

Install a Package:

```
sudo apt install nginx
```

Output: Installs nginx and its dependencies

3.

Remove a Package:

```
sudo apt remove nginx
```

Removes nginx, keeps config files

```
sudo apt purge nginx
```

Removes nginx and config files

4.

Clean Up Unused Dependencies:

```
sudo apt autoremove  
# Removes unneeded dependencies
```

5.

Troubleshooting

- **Broken dependencies:** Run `sudo apt update && sudo apt -f install` to fix.
- **Repository errors:** Check `/etc/apt/sources.list` for invalid URLs.
- **Permission denied:** Ensure `sudo` is used.
- **Package not found:** Verify package name or add a repository (`add-apt-repository`).

2. yum/dnf (Yellowdog Updater Modified/Dandified YUM)

Overview

- **Purpose:** High-level package managers for Red Hat-based systems (`yum` for older systems like CentOS 7, `dnf` for newer systems like Fedora or CentOS 8+). They manage `.rpm` packages and repositories.
- **Mechanics:** Queries repositories defined in `/etc/yum.repos.d/` (or `/etc/dnf/dnf.conf` for `dnf`) to handle package installation and updates.
- **Use Cases:**
 - Installing software on CentOS or Fedora.
 - Managing repository priorities.
 - Updating system packages.

Common Options (yum/dnf)

- `update` or `upgrade`: Update packages.
- `install [package]`: Install a package.
- `remove [package]`: Remove a package.
- `clean all`: Clear cached data.
- `--enablerepo=[repo]`: Enable a specific repository.
- `-y`: Assume "yes" to prompts.

Detailed Examples

Update Package Lists (dnf):

```
sudo dnf update
```

Output: Refreshes repository metadata and updates packages

1.

Install a Package (yum):

```
sudo yum install httpd
```

Output: Installs Apache web server and dependencies

2.

Remove a Package (dnf):

```
sudo dnf remove httpd
```

Output: Removes Apache and dependencies

3.

Enable a Specific Repository:

```
sudo dnf --enablerepo=epel install mypackage
```

Output: Installs mypackage from EPEL repository

4.

Troubleshooting

- **Dependency conflicts:** Use `dnf check` or `yum check` to diagnose.
- **Repository errors:** Verify `/etc/yum.repos.d/` files for correct URLs.
- **Slow performance:** Clear cache (`sudo dnf clean all` or `sudo yum clean all`).
- **Package not found:** Enable additional repositories (e.g., EPEL with `sudo yum install epel-release`).

3. dpkg (Debian Package Manager)

Overview

- **Purpose:** A low-level tool for installing, removing, and querying `.deb` package files on Debian-based systems.
- **Mechanics:** Directly manipulates individual `.deb` files, without handling dependencies or repositories (unlike `apt`).

- **Use Cases:**
 - Installing a downloaded `.deb` file.
 - Querying installed package details.
 - Fixing broken package installations.

Common Options

- `-i [file]`: Install a `.deb` file.
- `-r [package]`: Remove a package.
- `-P [package]`: Purge a package and its config files.
- `-l [pattern]`: List installed packages matching a pattern.
- `-s [package]`: Show package status.

Detailed Examples

Install a `.deb` File:

```
sudo dpkg -i package.deb  
# Output: Installs package.deb
```

1.

List Installed Packages:

```
dpkg -l | grep nginx  
# Output: Lists installed nginx packages
```

2.

Remove a Package:

```
sudo dpkg -r nginx  
# Removes nginx, keeps config files  
sudo dpkg -P nginx  
# Purges nginx and config files
```

3.

Fix Dependency Issues:

```
sudo dpkg -i package.deb  
# If errors occur due to dependencies:  
sudo apt -f install  
# Resolves missing dependencies
```

4.

Troubleshooting

- **Dependency errors:** Use `apt -f install` to fix after `dpkg -i`.
- **Package conflicts:** Remove conflicting packages (`dpkg -r`).
- **Broken packages:** Use `sudo dpkg --configure -a` to reconfigure.

4. rpm (RPM Package Manager)

Overview

- **Purpose:** A low-level tool for installing, removing, and querying `.rpm` package files on Red Hat-based systems.
- **Mechanics:** Directly handles `.rpm` files, without dependency resolution or repository management (unlike `yum/dnf`).
- **Use Cases:**
 - Installing a downloaded `.rpm` file.
 - Querying package metadata.
 - Verifying installed packages.

Common Options

- `-i [file]`: Install an `.rpm` file.
- `-e [package]`: Erase (remove) a package.
- `-qa`: Query all installed packages.
- `-qi [package]`: Show package information.
- `-qf [file]`: Find which package owns a file.

Detailed Examples

Install an .rpm File:

```
sudo rpm -i package.rpm  
# Output: Installs package.rpm
```

1.

List All Installed Packages:

```
rpm -qa | grep httpd  
# Output: Lists installed httpd packages
```

2.

Remove a Package:

```
sudo rpm -e httpd  
# Removes httpd
```

3.

Find Package Owning a File:

```
rpm -qf /usr/sbin/httpd  
# Output: e.g., httpd-2.4.37-1.el8
```

4.

Troubleshooting

- **Dependency errors:** Use `yum install` or `dnf install` to resolve dependencies after `rpm -i`.
- **Package already installed:** Use `--replacepks` with `rpm -i`.
- **Corrupted database:** Rebuild RPM database (`sudo rpm --rebuilddb`).

Practical Tips

- **Use high-level tools:** Prefer `apt` or `yum/dnf` over `dpkg/rpm` for dependency management.
- **Update regularly:** Run `apt update && apt upgrade` or `dnf update` to keep systems secure.
- **Check repositories:** Ensure `/etc/apt/sources.list` or `/etc/yum.repos.d/` are correctly configured.
- **Clean caches:** Use `apt autoclean` or `dnf clean all` to free disk space.
- **Verify packages:** Use `dpkg -l` or `rpm -qa` to list installed software.
- **Backup configs:** Copy `/etc/apt` or `/etc/yum.repos.d/` before changes.

Practice Tasks

1. Update package lists using `apt` or `dnf`.
2. Install `htop` using `apt` or `yum/dnf`.
3. List installed packages matching "nginx" using `dpkg` or `rpm`.
4. Remove `htop` using `apt` or `yum/dnf`.
5. Install a downloaded `.deb` or `.rpm` file using `dpkg` or `rpm`.

6. Check which package owns `/usr/bin/vim`.

Solution (Debian-based):

```
sudo apt update
sudo apt install htop
dpkg -l | grep nginx
sudo apt remove htop
sudo dpkg -i package.deb
sudo apt -f install # Fix dependencies
dpkg -S /usr/bin/vim
```

Solution (Red Hat-based):

```
sudo dnf update
sudo dnf install htop
rpm -qa | grep nginx
sudo dnf remove htop
sudo rpm -i package.rpm
sudo dnf install -y # Fix dependencies
rpm -qf /usr/bin/vim
```


Linux Firewall Management

This document provides an in-depth explanation of Linux commands and tools for managing firewalls: `ufw`, `iptables`, and `firewalld`. Each section includes a detailed description, how the tool works, common options, practical examples, use cases, and troubleshooting tips to help you configure and manage firewall rules effectively.

Overview of Firewall Management

- **Purpose:** Firewalls control incoming and outgoing network traffic based on predefined rules, enhancing system security by blocking unauthorized access and allowing legitimate traffic.
- **Mechanics:** Firewalls operate at the kernel level using `netfilter`, with tools like `ufw`, `iptables`, and `firewalld` providing user interfaces to configure rules.
- **Use Cases:**
 - Securing a web server by allowing only HTTP/HTTPS traffic.
 - Blocking specific IP addresses or ports.
 - Managing dynamic firewall rules for complex environments.
- **Key Concepts:**
 - **Rules:** Define actions (e.g., accept, drop) based on criteria (e.g., port, IP, protocol).
 - **Chains:** Categories for rules (e.g., INPUT, OUTPUT, FORWARD).
 - **Policies:** Default actions for unmatched traffic (e.g., DROP, ACCEPT).

1. ufw (Uncomplicated Firewall)

Overview

- **Purpose:** A user-friendly frontend for `iptables`, simplifying firewall configuration on Debian-based systems (e.g., Ubuntu).
- **Mechanics:** Translates high-level commands into `iptables` rules, managing the `netfilter` framework.
- **Use Cases:**
 - Quickly enabling/disabling firewall rules.
 - Allowing specific services (e.g., SSH, HTTP).
 - Blocking unwanted traffic for basic server security.

Common Options

- `enable`: Enable the firewall.
- `disable`: Disable the firewall.

- **status**: Show current rules and status.
- **allow [port/protocol]**: Allow traffic on a port or service.
- **deny [port/protocol]**: Block traffic on a port or service.
- **delete [rule]**: Remove a rule.
- **reset**: Reset to default (disables firewall, clears rules).

Detailed Examples

Enable ufw and Allow SSH:

```
sudo ufw enable
sudo ufw allow ssh
# Allows port 22/tcp
sudo ufw status
# Output: Status: active
#      To      Action    From
#      22/tcp    ALLOW    Anywhere
```

1.

Allow HTTP and HTTPS:

```
sudo ufw allow http
sudo ufw allow https
# Allows ports 80/tcp and 443/tcp
```

2.

Deny Specific Port:

```
sudo ufw deny 23/tcp
# Blocks telnet
```

3.

Allow Specific IP:

```
sudo ufw allow from 192.168.1.100
# Allows all traffic from 192.168.1.100
```

4.

Delete a Rule:

```
sudo ufw status numbered
```

```
# Output: [1] 22/tcp ALLOW Anywhere
sudo ufw delete 1
```

5.

Troubleshooting

- **Locked out (e.g., SSH):** Ensure `ufw allow ssh` is set before enabling.
- **Rules not applying:** Check status (`sudo ufw status`) and reload (`sudo ufw reload`).
- **Permission denied:** Use `sudo` for all `ufw` commands.
- **Conflicts with iptables:** Avoid mixing `ufw` and direct `iptables` rules.

2. iptables (IP Tables)

Overview

- **Purpose:** A low-level tool to configure `netfilter` firewall rules, offering fine-grained control over network traffic.
- **Mechanics:** Defines rules in chains (INPUT, OUTPUT, FORWARD) to accept, drop, or reject packets based on criteria like port, IP, or protocol.
- **Use Cases:**
 - Creating complex firewall rules.
 - Logging specific traffic.
 - Managing NAT (Network Address Translation).

Common Options

- `-A [chain]`: Append a rule to a chain.
- `-D [chain] [rule]`: Delete a rule.
- `-L [chain]`: List rules in a chain.
- `-F [chain]`: Flush (clear) rules in a chain.
- `-P [chain] [policy]`: Set default policy (e.g., DROP, ACCEPT).
- `-p [protocol]`: Specify protocol (e.g., tcp, udp).
- `--dport [port]`: Destination port.
- `-s [source]`: Source IP.
- `-j [target]`: Action (e.g., ACCEPT, DROP, REJECT).

Detailed Examples

List All Rules:

```
sudo iptables -L -v -n
# Output: Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
#      pkts bytes target prot opt in out source destination
```

1.

Allow SSH Traffic:

```
sudo iptables -A INPUT -p tcp --dport 22 -j ACCEPT
# Allows incoming SSH
```

2.

Block Specific IP:

```
sudo iptables -A INPUT -s 192.168.1.100 -j DROP
# Drops all traffic from 192.168.1.100
```

3.

Set Default Policy to DROP:

```
sudo iptables -P INPUT DROP
sudo iptables -A INPUT -p tcp --dport 80 -j ACCEPT
# Drops all incoming traffic except HTTP
```

4.

Save Rules:

```
sudo iptables-save > /etc/iptables/rules.v4
# Persist rules (Debian-based, requires iptables-persistent package)
```

5.

Troubleshooting

- **Rules not persistent:** Use `iptables-save` and a package like `iptables-persistent`.
- **Locked out:** Ensure critical ports (e.g., SSH) are allowed before setting DROP policy.
- **Complex rules:** Test rules incrementally and use `-L -v -n` to verify.
- **Conflicts with ufw/firewalld:** Use only one firewall tool to avoid conflicts.

3. firewalld (Firewall Daemon)

Overview

- **Purpose:** A dynamic firewall management tool for Red Hat-based systems (e.g., CentOS, Fedora), providing a high-level interface for `netfilter`.
- **Mechanics:** Uses zones (e.g., public, trusted) and services to manage rules, with dynamic updates via D-Bus.
- **Use Cases:**
 - Managing firewall rules on servers with changing network conditions.
 - Supporting NetworkManager integration.
 - Configuring complex firewall policies with zones.

Common Options

- `--get-zones`: List available zones.
- `--get-active-zones`: Show currently active zones.
- `--add-service=[service]`: Allow a service (e.g., `http`, `ssh`).
- `--add-port=[port/protocol]`: Allow a specific port.
- `--remove-service=[service]`: Remove a service rule.
- `--permanent`: Make changes persistent across reboots.
- `--reload`: Apply changes without interrupting connections.

Detailed Examples

Check Active Zones:

```
sudo firewall-cmd --get-active-zones
# Output: public
#       interfaces: eth0
```

1.

Allow HTTP Service:

```
sudo firewall-cmd --add-service=http
sudo firewall-cmd --permanent --add-service=http
sudo firewall-cmd --reload
# Output: success
```

2.

Allow Specific Port:

```
sudo firewall-cmd --add-port=8080/tcp --permanent  
sudo firewall-cmd --reload
```

3.

List Rules in a Zone:

```
sudo firewall-cmd --list-all  
# Output: public (active)  
#    services: ssh http  
#    ports: 8080/tcp
```

4.

Block an IP:

```
sudo firewall-cmd --add-source=192.168.1.100 --zone=drop --permanent  
sudo firewall-cmd --reload
```

5.

Troubleshooting

- **Rules not applied:** Use `--permanent` and `--reload` for persistence.
- **Service not found:** Check available services (`firewall-cmd --get-services`).
- **Locked out:** Ensure `ssh` or management ports are allowed in the active zone.
- **Conflicts with iptables/ufw:** Disable other firewall tools (`systemctl stop ufw`).

Practical Tips

- **Choose one tool:** Use `ufw` (Ubuntu), `firewalld` (Red Hat), or `iptables` exclusively to avoid conflicts.
- **Test rules:** Apply rules temporarily before making them permanent.
- **Backup rules:** Save `iptables` rules (``iptables`

Linux SSH and Remote Access

This document provides an in-depth explanation of Linux tools and configurations for secure remote access: `ssh`, `scp`, `sshd_config`, and key-based authentication. Each section includes a detailed description, how the tool or configuration works, common options, practical examples, use cases, and troubleshooting tips to help you manage remote connections securely.

Overview of SSH and Remote Access

- **Purpose:** Secure Shell (SSH) enables secure remote access to systems, allowing command execution, file transfers, and server management over encrypted connections.
- **Mechanics:** SSH uses client-server architecture with public-key cryptography for authentication and encryption. The SSH daemon (`sshd`) runs on the server, and clients use `ssh` or `scp` to connect.
- **Use Cases:**
 - Managing remote servers (e.g., web or database servers).
 - Securely transferring files between systems.
 - Setting up passwordless authentication for automation.
- **Key Files:**
 - `/etc/ssh/sshd_config`: SSH server configuration.
 - `~/.ssh/`: User directory for SSH keys and configs.
 - `/var/log/auth.log` or `/var/log/secure`: Logs for SSH authentication events.

1. ssh (Secure Shell Client)

Overview

- **Purpose:** Connects to a remote system to execute commands or start an interactive shell securely.
- **Mechanics:** Establishes an encrypted connection to the SSH daemon on the remote host, authenticating via password or key.
- **Use Cases:**
 - Running commands on a remote server.
 - Accessing a remote shell for administration.
 - Tunneling network traffic.

Common Options

- `-p [port]`: Specify SSH port (default: 22).
- `-i [keyfile]`: Use a specific private key for authentication.

- `-X`: Enable X11 forwarding for GUI applications.
- `-v`: Verbose output for debugging.
- `[user@host]`: Specify remote user and hostname/IP.

Detailed Examples

Connect to a Remote Server:

```
ssh user1@192.168.1.100
# Prompts for password, opens remote shell
```

1.

Connect with Custom Port:

```
ssh -p 2222 user1@192.168.1.100
# Connects to SSH server on port 2222
```

2.

Run a Single Command:

```
ssh user1@192.168.1.100 "uptime"
# Output: 22:43:45 up 5 days, 3:12, 2 users, load average: 0.10, 0.15, 0.20
```

3.

Enable X11 Forwarding:

```
ssh -X user1@192.168.1.100
# Allows running GUI apps (e.g., xterm) remotely
```

4.

Troubleshooting

- **Connection refused:** Ensure `sshd` is running (`sudo systemctl status sshd`) and the port is open (`ss -tln`).
- **Permission denied:** Verify username, password, or key; check `/var/log/auth.log`.
- **Timeout:** Check firewall rules (`ufw status` or `firewall-cmd --list-all`) and network connectivity (`ping`).

2. scp (Secure Copy)

Overview

- **Purpose:** Securely transfers files between local and remote systems or between two remote systems using SSH.
- **Mechanics:** Uses SSH for encryption, copying files over a secure channel.
- **Use Cases:**
 - Uploading configuration files to a server.
 - Downloading logs from a remote system.
 - Transferring files between servers.

Common Options

- **-P [port]:** Specify SSH port.
- **-r:** Recursively copy directories.
- **-i [keyfile]:** Use a specific private key.
- **-v:** Verbose output for debugging.

Detailed Examples

Copy a File to a Remote Server:

```
scp localfile.txt user1@192.168.1.100:/home/user1/  
# Copies localfile.txt to remote /home/user1/
```

1.

Copy a Directory Recursively:

```
scp -r /local/dir user1@192.168.1.100:/home/user1/  
# Copies entire directory
```

2.

Copy from Remote to Local:

```
scp user1@192.168.1.100:/home/user1/remotefile.txt .  
# Copies remotefile.txt to current directory
```

3.

Copy with Custom Port:

```
scp -P 2222 localfile.txt user1@192.168.1.100:/home/user1/
```

4.

Troubleshooting

- **Permission denied:** Check file permissions and SSH authentication.
- **Connection issues:** Verify SSH connectivity (`ssh user1@192.168.1.100`) and port.
- **Slow transfers:** Use `-v` to debug or check network bandwidth.

3. sshd_config (SSH Daemon Configuration)

Overview

- **Purpose:** Configures the SSH server (`sshd`) behavior, including port, authentication methods, and access controls.
- **Mechanics:** Located at `/etc/ssh/sshd_config`, defines settings like port, allowed users, and key-based authentication.
- **Use Cases:**
 - Changing the default SSH port for security.
 - Disabling password authentication.
 - Restricting access to specific users or groups.

Common Directives

- `Port [number]`: Set SSH port (default: 22).
- `PermitRootLogin [yes|no]`: Allow/disallow root login.
- `PasswordAuthentication [yes|no]`: Enable/disable password login.
- `AllowUsers [user1 user2]`: Restrict SSH to specific users.
- `PubkeyAuthentication [yes|no]`: Enable/disable key-based authentication.

Detailed Examples

Change SSH Port:

```
sudo nano /etc/ssh/sshd_config
# Change: Port 22
# To:    Port 2222
# Save and exit
sudo systemctl restart sshd
```

1.

Disable Root Login:

```
sudo nano /etc/ssh/sshd_config
# Set: PermitRootLogin no
```

```
sudo systemctl restart sshd
```

2.

Restrict to Specific Users:

```
sudo nano /etc/ssh/sshd_config  
# Add: AllowUsers user1 user2  
sudo systemctl restart sshd
```

3.

Disable Password Authentication:

```
sudo nano /etc/ssh/sshd_config  
# Set: PasswordAuthentication no  
sudo systemctl restart sshd
```

4.

Troubleshooting

- **Syntax errors:** Test config with `sudo sshd -t` before restarting.
- **Locked out:** Always test changes with an active session; ensure `AllowUsers` includes your user.
- **Service not restarting:** Check logs (`journalctl -u sshd`) for errors.
- **Firewall issues:** Update firewall rules for new port (`ufw allow 2222` or `firewall-cmd --add-port=2222/tcp`).

4. Key-Based Authentication

Overview

- **Purpose:** Enables passwordless SSH login using public-private key pairs, improving security and automation.
- **Mechanics:** A private key is stored on the client, and the corresponding public key is added to `~/.ssh/authorized_keys` on the server.
- **Use Cases:**
 - Automating scripts requiring SSH access.
 - Enhancing security by disabling passwords.
 - Simplifying logins for frequent access.

Steps to Set Up

Generate Key Pair on Client:

```
ssh-keygen -t rsa -b 4096  
# Press Enter for defaults (~/.ssh/id_rsa)
```

1.

Copy Public Key to Server:

```
ssh-copy-id user1@192.168.1.100  
# Adds public key to ~user1/.ssh/authorized_keys on server
```

2.

Test Key-Based Login:

```
ssh user1@192.168.1.100  
# Should log in without password
```

3.

Disable Password Authentication (Optional):

```
sudo nano /etc/ssh/sshd_config  
# Set: PasswordAuthentication no  
sudo systemctl restart sshd
```

4.

Detailed Example

Full Setup:

```
# Generate key pair  
ssh-keygen -t rsa -b 4096  
# Copy key to server  
ssh-copy-id -i ~/.ssh/id_rsa.pub user1@192.168.1.100  
# Test login  
ssh user1@192.168.1.100  
# Verify authorized_keys  
ssh user1@192.168.1.100 "cat ~/.ssh/authorized_keys"
```

•

Troubleshooting

- **Permission denied:** Ensure `~/.ssh` is 700 and `authorized_keys` is 600 (`chmod 700 ~/.ssh; chmod 600 ~/.ssh/authorized_keys`).
- **Key rejected:** Check server logs (`sudo journalctl -u sshd`) and verify key in `authorized_keys`.
- **SELinux issues:** On Red Hat-based systems, restore context (`sudo restorecon -R -v ~/.ssh`).
- **Firewall blocks:** Ensure SSH port is open (`ufw status` or `firewall-cmd --list-all`).

Practical Tips

- **Secure SSH:** Change default port, disable root login, and use key-based authentication.
- **Backup keys:** Store `~/.ssh/id_rsa` securely; losing it prevents access.
- **Monitor logs:** Check `/var/log/auth.log` or `/var/log/secure` for SSH issues.
- **Use `scp` for automation:** Combine with key-based authentication for scripts.
- **Test config changes:** Use `sshd -t` and keep an active session to avoid lockouts.
- **Combine with `tmux`:** Use `tmux` for persistent remote sessions.

Practice Tasks

1. Connect to a remote server (`user1@192.168.1.100`) using `ssh`.
2. Copy a local file (`data.txt`) to the remote server using `scp`.
3. Change the SSH port to 2222 in `sshd_config` and restart `sshd`.
4. Set up key-based authentication for `user1@192.168.1.100`.
5. Disable password authentication on the server.
6. Verify SSH login without a password.

Solution:

```
ssh user1@192.168.1.100
scp data.txt user1@192.168.1.100:/home/user1/
sudo nano /etc/ssh/sshd_config
# Set: Port 2222
sudo systemctl restart sshd
ssh-keygen -t rsa -b 4096
ssh-copy-id -i ~/.ssh/id_rsa.pub user1@192.168.1.100
sudo nano /etc/ssh/sshd_config
# Set: PasswordAuthentication no
sudo systemctl restart sshd
ssh -p 2222 user1@192.168.1.100
```

Linux System Monitoring and Performance

This document provides an in-depth explanation of Linux commands for system monitoring and performance analysis: `sar`, `vmstat`, `iostat`, and `free`. Each section includes a detailed description, how the tool works, common options, practical examples, use cases, and troubleshooting tips to help you monitor system resources effectively.

Overview of System Monitoring

- **Purpose:** System monitoring tools track resource usage (CPU, memory, disk, I/O) to diagnose performance issues, optimize system performance, and ensure stability.
- **Mechanics:** These tools query kernel data (e.g., `/proc`, `/sys`) to provide real-time or historical metrics on system performance.
- **Use Cases:**
 - Identifying CPU or memory bottlenecks.
 - Monitoring disk I/O performance.
 - Analyzing system load for capacity planning.
- **Key Files:**
 - `/proc/stat`: CPU and system statistics.
 - `/proc/meminfo`: Memory usage details.
 - `/proc/diskstats`: Disk I/O statistics.
 - `/var/log/sysstat/`: Historical data for `sar`.

1. sar (System Activity Reporter)

Overview

- **Purpose:** Collects, reports, and saves system activity statistics (CPU, memory, I/O, etc.) over time, part of the `sysstat` package.
- **Mechanics:** Gathers data via cron (`/etc/cron.d/sysstat`) and stores it in `/var/log/sysstat/saXX` for historical analysis.
- **Use Cases:**
 - Analyzing historical CPU usage trends.
 - Troubleshooting intermittent performance issues.
 - Generating performance reports.

Common Options

- **-u**: CPU usage (default).
- **-r**: Memory usage.
- **-b**: I/O and transfer rate statistics.
- **-d**: Disk activity (use **-p** for device names).
- **-A**: All statistics.
- **-f [file]**: Read from a specific **sa** file.
- **-s [start] -e [end]**: Specify time range (e.g., **10:00:00**).

Detailed Examples

View Current CPU Usage:

```
sar -u 1 5
# Output: CPU usage every 1 second, 5 times
#    22:45:01  CPU %user %nice %system %iowait %steal %idle
#    22:45:02  all 5.00 0.00  2.00   0.50  0.00 92.50
```

1.

Check Historical Memory Usage:

```
sar -r -f /var/log/sysstat/sa27
# Output: Memory stats for June 27
#    10:00:01  kbmfree kbmemused %memused kbbuffers kbcached
#    10:00:02  500000  1500000  75.00   100000  200000
```

2.

Disk I/O for a Specific Device:

```
sar -d -p 1 5
# Output: Disk stats every 1 second, 5 times
#    22:45:01  DEV    tps  kB/s  kB/s
#    22:45:02  sda    10.0  50.0  20.0
```

3.

All Stats for a Time Range:

```
sar -A -s 10:00:00 -e 12:00:00
# Output: All metrics for specified time range
```

4.

Troubleshooting

- **Command not found:** Install `sysstat` (`sudo apt install sysstat` or `sudo dnf install sysstat`).
- **No data in files:** Enable data collection in `/etc/sysstat/sysstat` (`ENABLED="true"`) and restart (`sudo systemctl restart sysstat`).
- **Missing historical data:** Check `/var/log/sysstat/` and cron job (`/etc/cron.d/sysstat`).
- **Permission denied:** Use `sudo` for restricted files.

2. vmstat (Virtual Memory Statistics)

Overview

- **Purpose:** Reports system performance metrics, including CPU, memory, swap, I/O, and system activity.
- **Mechanics:** Reads `/proc/stat`, `/proc/meminfo`, and `/proc/vmstat` for real-time data.
- **Use Cases:**
 - Monitoring memory and swap usage.
 - Detecting I/O bottlenecks.
 - Checking system load during high activity.

Common Options

- `-s`: Summary of memory and CPU stats.
- `-d`: Disk statistics.
- `-a`: Show active/inactive memory.
- `[delay [count]]`: Update every `delay` seconds, `count` times.

Detailed Examples

Monitor System Activity:

```
vmstat 2 5
# Output: Updates every 2 seconds, 5 times
#   procs  memory   swap   io   system  cpu
#   r  b  swpd  free  buff  cache  si  so  bi  bo  in  cs  us  sy  id  wa
#   1  0  0    500M  100M  1G    0  0   10 20 100 200 5  2  90 3
```

1.

Memory Summary:

```
vmstat -s
```



```
# Output:      2000000 K total memory
#             500000 K used memory
#             300000 K active memory
```

2.

Disk Statistics:

```
vmstat -d
# Output: disk reads merged sectors ms writes merged sectors ms
#      sda 1000 500 8000 200 2000 1000 16000 400
```

3.

Troubleshooting

- **No output:** Ensure `procps` is installed (`sudo apt install procps`).
- **High swap usage (`si/so`):** Check memory pressure (`free -m`) and add RAM if needed.
- **Unclear metrics:** Refer to `man vmstat` for column explanations (e.g., `si` = swap in).

3. iostat (Input/Output Statistics)

Overview

- **Purpose:** Reports CPU and disk I/O statistics, part of the `sysstat` package.
- **Mechanics:** Reads `/proc/diskstats` and `/proc/stat` for disk and CPU performance data.
- **Use Cases:**
 - Identifying disk performance bottlenecks.
 - Monitoring CPU utilization alongside I/O.
 - Analyzing storage device performance.

Common Options

- `-c`: CPU statistics only.
- `-d`: Device (disk) statistics only.
- `-x`: Extended statistics (e.g., %util).
- `-p [device]`: Show stats for a specific device.
- `[interval [count]]`: Update every `interval` seconds, `count` times.

Detailed Examples

Monitor CPU and Disk I/O:

```
iostat 2 5
# Output: CPU and device stats every 2 seconds, 5 times
#      avg-cpu:  %user  %nice %system %iowait %steal %idle
#              5.00  0.00  2.00  0.50  0.00  92.50
#      Device tps kB_read/s kB_wrtn/s
#      sda   10.0 50.0    20.0
```

1.

Extended Disk Stats:

```
iostat -x
# Output: Device tps kB_read/s kB_wrtn/s %util
#      sda   10.0 50.0    20.0    5.0
```

2.

Specific Device:

```
iostat -p sda
# Output: Stats for sda and its partitions
```

3.

Troubleshooting

- **Command not found:** Install `sysstat` (`sudo apt install sysstat`).
- **High %iowait:** Indicates disk bottleneck; check device utilization (`-x`).
- **No device stats:** Ensure disks are active (`lsblk`).

4. free (Memory Usage)

Overview

- **Purpose:** Displays memory and swap usage, including total, used, and free amounts.
- **Mechanics:** Reads `/proc/meminfo` for memory statistics.
- **Use Cases:**
 - Checking available RAM for applications.
 - Monitoring swap usage.
 - Diagnosing memory shortages.

Common Options

- **-h**: Human-readable format (e.g., GB, MB).
- **-m**: Display in megabytes.
- **-s [seconds]**: Continuous updates every **seconds**.
- **-t**: Show total for memory and swap.

Detailed Examples

View Memory Usage:

```
free -h
# Output:      total used free shared buff/cache available
#   Mem:      2.0G 500M 1.0G 100M  400M    1.2G
#   Swap:      1.0G  0B 1.0G
```

1.

Continuous Monitoring:

```
free -h -s 2
# Updates every 2 seconds
```

2.

Show Totals:

```
free -h -t
# Output: Includes a total line for Mem + Swap
```

3.

Troubleshooting

- **High swap usage**: Indicates low RAM; consider adding memory or closing apps.
- **Low available memory**: Check **buff/cache** (can be reclaimed) vs. **used**.
- **Inaccurate stats**: Cross-check with **vmstat -s** or **top**.

Practical Tips

- **Use **sar** for historical data**: Ideal for long-term analysis.
- **Combine tools**: Use **vmstat** for quick snapshots, **iostat** for disk details, and **free** for memory.
- **Monitor regularly**: Set up **sar** cron jobs for continuous data collection.
- **Check logs**: Use **/var/log/syslog** or **journalctl** for related errors.

- **Automate alerts:** Script checks for high CPU/memory usage with `sar` or `vmstat`.
- **Install sysstat:** Ensure `sysstat` is installed for `sar` and `iostat`.

Practice Tasks

1. Check current CPU usage with `sar` for 3 updates every 2 seconds.
2. Display memory statistics with `vmstat`.
3. Monitor disk I/O for `sda` with `iostat`.
4. View memory and swap usage in human-readable format with `free`.
5. Check historical CPU usage for today with `sar`.
6. Combine `vmstat` and `free` to monitor memory usage every 5 seconds.

Solution:

```
sar -u 2 3
vmstat -s
iostat -p sda
free -h
sar -u -f /var/log/sysstat/sa27
vmstat 5 & free -h -s 5
# Press Ctrl+C to stop
```